

The Making of Dominant Vehicle Currencies: Evidence from DeFi

August 20, 2026

Abstract

We use 472 million swaps to observe vehicle currencies in transaction-level routes on decentralised exchanges. Between matched January–June periods in 2024 and 2026, stablecoins’ share among native and stable vehicles rises from 16.9% to 42.3% by route count. Positive and negative substitutions offset within endpoint pairs active in both periods. The aggregate increase comes from activity shifting toward continuing pairs that already rely more on stablecoins and from pairs entering or leaving the trading network. Vehicle choice at pair entry predicts later choice, and stablecoin route share rises with relative deposited capital. Trading relationships and local liquidity jointly shape vehicle dominance.

1 Introduction

Every cross-border payment between two currencies needs an institution willing to receive one and deliver the other. When a payment provider neither participates in the destination payment system nor holds the destination currency, the BIS Project Nexus design assigns that task to a third-party foreign-exchange provider ([Bank for International Settlements Innovation Hub, 2024](#)). Project Rialto goes one step further: its experimental cross-border system uses a third currency when direct exchange is unavailable ([Bank for International Settlements Innovation Hub, 2025](#)). The practical question is which currency connects a thin trading relationship and what liquidity makes that connection viable.

The same problem underlies the dominance of the U.S. dollar. The dollar appears on more than 85% of foreign-exchange transactions, far more often than the United States’ share of cross-border economic activity would suggest ([Somogyi, 2026](#)). Its use as a vehicle between other currencies concentrates trading in dollar markets; deeper markets lower trading costs and reinforce the dollar’s intermediary role ([Krugman, 1980](#); [Kiyotaki and Wright, 1989](#); [Flandreau and Jobst, 2009](#)). This feedback explains persistence after a vehicle has become dominant. It says less about how a challenger gains share when new trading relationships and new market technologies appear.

Vehicle dominance depends on both the intermediary selected within a source–destination pair and the amount traded across different source–destination pairs. Aggregate dominance can therefore change even when the vehicle used inside every continuing pair remains unchanged. Measuring

these two channels requires transactions that attach each intermediary to the exchange between the final source and destination.

Standard foreign-exchange records report turnover in each currency pair without linking the legs of the customer’s exchange.¹ A yen–euro exchange routed through dollars appears as separate yen–dollar and dollar–euro trades. [Somogyi \(2026\)](#) infers vehicle activity from currency-pair volume responses around non-overlapping holidays. In decentralised exchange, by contrast, the pool calls inside one blockchain transaction preserve their execution order. We use these records to observe which asset connects the source to the destination.

We call the ordered source and destination the *endpoint pair*, and each pool trade a *leg*. A route may contain one leg, several sequential legs, or a split across pools. We reconstruct connected routes from 472 million swaps on eight Ethereum automated market maker (AMM) deployments between February 2020 and June 2026. The resulting token sequences identify the observed endpoint pair, the intermediary, and whether the route spans exchanges. They reveal vehicle use directly, avoiding inference from turnover in separate legs.

Stablecoins, including USDC, USDT, DAI, and other pegged tokens, become substantially more prominent vehicles during the sample. We compare the same January-through-June dates in 2024 and 2026 because the 2026 data end on June 30. Stablecoin dominance rises from 16.9% to 42.3% by route count and from 32.7% to 76.5% by dollar-weighted volume. Within endpoint pairs active in both periods, positive and negative stable-for-native substitutions offset in aggregate. The increase instead comes from trading shifting toward continuing pairs that already rely more on stablecoins and from pairs observed in only one comparison period.

The rotation is not a uniform move into every stablecoin. Native assets regain share in 2023 and 2024 before stablecoins recover the dollar-weighted lead in 2025. USDT’s intermediary use rises especially sharply relative to its appearance at route endpoints. The stablecoin gain appears on routes confined to one exchange and on routes spanning several exchanges.

Entry does more than change the set of pairs. The vehicle used when an endpoint pair first appears strongly predicts the vehicle used over the next 120 days, including after removing pairs with WETH at an endpoint and controlling for entry size and route structure. Local liquidity then predicts which vehicle carries a trade once both alternatives are present. During the first month after a persistent stablecoin bridge forms, stablecoin depth on the shallower leg is below one-tenth of WETH depth on most pair-days. Stablecoins carry 53.0% of routes when that depth is at least WETH depth and 69.9% when it is at least twice WETH depth. These descriptive associations locate the vehicle role where pair formation and deposited capital meet; they do not identify an exogenous liquidity-supply effect.

Directly observed routes change the interpretation of vehicle-currency dominance. A large aggregate rotation can arise because the trading network expands around a challenger, even when continuing pairs show little net substitution. This formation channel complements the holiday-based inference of [Somogyi \(2026\)](#) and the historical network evidence of [Flandreau and Jobst \(2009\)](#).

¹For example, see the [BIS 2025 Triennial Survey reporting guidelines and templates](#).

The persistence and depth results connect that channel to classical theories in which vehicle use and liquidity reinforce one another, and to fragmented-market research in which liquidity supply determines the trading opportunities available to investors (Krugman, 1980; Chen and Duffie, 2021; Lehar and Parlour, 2025). A vehicle currency is made when it wins existing trades and when it becomes the liquid bridge inherited by new ones.

2 Setting and data

2.1 Routes and notation

Decentralised exchanges distribute liquidity across smart-contract pools that routers can combine atomically. A trader may use one pool, divide an order across several pools, or trade through an intermediary asset. Because the pool calls share a transaction, the intermediary sequence is observable.² A route through an intermediary is unavoidable when no direct pool connects the assets; when direct and indirect pools coexist, the router compares their fees, price impact, gas, and venue access. Lehar and Parlour (2025) describe the constant-product design and the economics of liquidity supply in these markets.

Blockchain logs record each pool call in execution order. We recover a route by following the directed token flows within a transaction. Calls with a common source and destination form a split order; a token received in one call and spent in a later call joins the calls into a sequential route. The same rule separates unrelated conversions bundled in one transaction. Throughout, for a reconstructed sequence $A \rightarrow B \rightarrow C$, we call (A, C) the ordered *endpoint pair* and each pool trade, $A \rightarrow B$ or $B \rightarrow C$, a *leg*. The order in (A, C) records the exchange direction, while the route records the intervening sequence $A \rightarrow B \rightarrow C$. Each leg trades in a liquidity pool; an endpoint pair need not have a direct pool. These endpoints belong to the connected pool component and need not coincide with the endpoints of a broader user instruction. We use *pair* as shorthand for endpoint pair only where no pool-level pair is under discussion.

A transaction on January 10, 2026 illustrates both the information in the logs and its limit. A public explorer describes a 460,000 PYUSD–USDe swap executed through 0x.³ The user’s instruction begins with PYUSD, but the connected pool sequence begins when 459,929.81 USDC enters Fluid. Fluid returns 460,331.05 USDT, which then enters Uniswap v4 and produces 460,104.28 USDe. The observed endpoint pair is therefore USDC–USDe, the two legs are USDC→USDT and USDT→USDe, and USDT is the vehicle. The earlier PYUSD–USDC conversion occurs outside the connected pool sequence, so the data do not show whether USDC was another vehicle or inventory supplied by the executor. We apply the connected-component rule to every transaction and do not infer an unobserved leg from nearby transfers. This choice may understate the intermediary role of assets such as USDC when a broader instruction is funded outside the observed pool sequence.

The sequence directly identifies intermediary use. Conditional on reconstructed pool states

²Ethereum records a transaction’s ordered logs in its receipt; see the official [Ethereum JSON-RPC documentation](#).

³Transaction [0xbda1d07f...9bf67ad9](#).

immediately before execution, it also defines a same-size cost comparison between the observed path and available alternatives. We introduce notation for the route here and develop the cost comparison in [Section 4](#).

Write $\mathcal{A}_t$ for the assets tradable on day t and $\mathcal{P}_t$ for the pools live on that day. For a route r , let $e(r)$ record its block, transaction position, and first pool-event position. The notation $e(r)^-$ refers to the state immediately before that transaction begins; a daily closing state generally occurs later. A route carries an ordered pool sequence $(p_1, \dots, p_{m(r)})$ and token sequence $(k_0, \dots, k_{m(r)})$, where $k_0 = i$ is the input and $k_{m(r)} = o$ is the output. Thus (i, o) is the endpoint pair and each (k_{h-1}, k_h) is one leg. The dollar input notional q corresponds to $\Delta_i(q, e) = q/P_{i,e}$ units of the input token. Let χ_{p,e^-} denote pool p 's state immediately before execution and let $Q_p(i \rightarrow j, \Delta_i; \chi_{p,e^-})$ denote the units of token j returned for an exact input of Δ_i units of token i . Composing these pool quotes gives the route's token output and dollar value,

$$\begin{aligned} \Delta_{k_h} &= Q_{p_h}(k_{h-1} \rightarrow k_h, \Delta_{k_{h-1}}; \chi_{p_h, e(r)^-}), & h &= 1, \dots, m(r), \\ Q_r(q, e(r)) &= \Delta_o, & O_r(q, e(r)) &= P_{o, e(r)} Q_r(q, e(r)). \end{aligned} \tag{1}$$

The venue-specific forms of Q_p appear in Appendix B. Pool invariants take token quantities as arguments, which is why the dollar notional q and the token input Δ_i remain distinct. Equation (1) describes an executed route and, once market state has been reconstructed, a hypothetical route.

Intermediary use is one of several currency roles. Invoicing concerns the currency named in trade contracts (Gopinath et al., 2020; Amiti et al., 2022; Mukhin, 2022); funding concerns the denomination of liabilities (Gopinath and Stein, 2021; Eren and Malamud, 2022); and dealer turnover records trading activity (Somogyi, 2026). These roles can reinforce one another. The path in Equation (1) identifies the asset used between the currency sold and the currency bought.

This route setting makes the classical volume–cost mechanism observable at transaction level. Greater use of a vehicle expands trading in the two markets joined through it; deeper markets can then lower the cost of using the same vehicle again (Krugman, 1980). In Kiyotaki and Wright (1989), traders accept an asset for indirect exchange when its marketability compensates for its storage cost. Dowd and Greenaway (1993) emphasise the fixed costs of learning, accounting, and quoting in a new currency. Liquidity placement, venue access, and the routes available to software provide corresponding coordination margins on decentralised exchanges.

2.2 Measuring vehicle dominance

An exact two-leg route consists of two atomic trades, one on either side of a unique intermediary. This gives the primary measure a simple unit: each ultimate trade contributes one intermediary observation. A supplementary multiple-intermediary measure includes routes of any length and counts each intermediary use separately. For route r , let $\mathcal{M}(r)$ denote the set of assets between the ultimate source and destination. We record the intermediary status of asset k as

$$Z_{r,k} = \mathbf{1}\{k \in \mathcal{M}(r)\}.$$

Let $a(k)$ assign each token to an economic asset type. The native group consolidates ETH and its one-for-one wrapped representation, WETH, after route reconstruction. The stable group con-

tains fiat-reserve, on-chain-collateralised, synthetic, and non-dollar stable currencies. For $g \in \{\text{native, stable}\}$, define $Z_{r,g} = \sum_{k:a(k)=g} Z_{r,k}$; on an exact two-leg route, this sum is an indicator.

This grouping preserves a token-level estimand. WETH is the one-for-one token representation that allows ETH to interact with ERC-20 pool contracts; wrapping changes the technical form, not the underlying settlement exposure. We therefore reconstruct routes at the token level and consolidate ETH and WETH only when assigning the economic asset used as a vehicle. Dollar-pegged tokens remain distinct instruments because their issuers, redemption arrangements, risks, and leg-level liquidity differ. USDT in the USDC–USDT–USDe example is therefore an intermediary token even though all three instruments target the same unit of account. The aggregate stable group measures dominance by a family of stable tokens, while the token-level results retain issuer-specific competition.

We begin with the competition between native and stable intermediaries. Let $\mathcal{R}_t^{(2),NS}$ denote exact two-leg routes on day t whose intermediary belongs to either group, and let $\mathcal{R}_t^{(2),NS,V}$ retain routes for which the source, intermediary, and destination values agree within 20%. The equally weighted and dollar-weighted group shares are

$$S_{g,t}^N = \frac{\sum_{r \in \mathcal{R}_t^{(2),NS}} Z_{r,g}}{|\mathcal{R}_t^{(2),NS}|}, \quad S_{g,t}^V = \frac{\sum_{r \in \mathcal{R}_t^{(2),NS,V}} \nu_r Z_{r,g}}{\sum_{r \in \mathcal{R}_t^{(2),NS,V}} \nu_r}, \quad (2)$$

where ν_r is the arithmetic mean of the route’s total dollar value at its observed source and destination. The 20% rule also requires every intermediary’s incoming and outgoing dollar values to have a ratio between 0.8 and 1.2. Thus $S_{\text{stable},t}^N$ is stablecoins’ aggregate share among native and stable intermediaries, and $S_{\text{stable},t}^V$ is the corresponding dollar-weighted share. The two groups exhaust this denominator. Share levels are reported in percent; the unit for absolute share changes is one percentage point (pp).

The frequency with which an asset appears at the beginning or end of an observed pool route provides a benchmark for its intermediary use. This comparison uses a broader set of currencies $\mathcal{C}$: native, staked-native, stable, and imported currencies, with unclassified assets excluded. We retain every complete route that does not return to its starting currency. Direct routes contribute to the source and destination counts, and a longer route contributes one observation for each intermediary it uses. Let $n_{k,t}^I$ count intermediary uses and $n_{k,t}^E$ count source or destination appearances for $k \in \mathcal{C}$. The two shares are

$$I_{k,t}^N = \frac{n_{k,t}^I}{\sum_{j \in \mathcal{C}} n_{j,t}^I}, \quad E_{k,t}^N = \frac{n_{k,t}^E}{\sum_{j \in \mathcal{C}} n_{j,t}^E},$$

and the corresponding share gap and excess-use ratio are

$$\Delta_{k,t}^N = I_{k,t}^N - E_{k,t}^N, \quad X_{k,t}^N = \frac{I_{k,t}^N}{E_{k,t}^N} \quad \text{when } E_{k,t}^N > 0. \quad (3)$$

A positive $\Delta_{k,t}^N$ means that the asset’s share of intermediary use exceeds its share of source and

destination appearances. We interpret $X_{k,t}^N$ as intermediation intensity relative to endpoint demand: zero means the asset is not used as an intermediary, one means its two shares are equal, and a value above one means it is overrepresented as a vehicle. The ratio is not a ranking of aggregate vehicle dominance. Dollar-weighted measures replace counts with route values while retaining the same currencies and the 20% agreement rule. Because this role-specialisation measure compares each currency’s intermediary and endpoint uses, its denominator differs from the native-or-stable denominator in Equation (2). The latter remains the measure of aggregate rotation.

Source and destination assets must differ in the analysis sample because vehicle exchange connects two distinct endpoint assets. We classify routes that return to their starting asset as round trips and remove them. Such paths may contain cyclic arbitrage or other self-returning strategies: their economic object is a trading cycle, not an exchange from one endpoint asset into another, although trader intent remains unobserved. [Daian et al. \(2020\)](#) show that several arbitrage trades can be bundled atomically. On the median of 79 sampled days, round trips account for 12.7% of multi-leg routes by count and 21.7% by value. Appendix A reports their distribution across the sampled dates.

Within this route sample, we classify assets by their economic exposure: native, staked native, stable, imported, and other. The native category contains the platform’s settlement asset. We treat ETH and its one-for-one wrapped representation, WETH, as one economic asset after reconstructing the route. This consolidation is an economic classification, not an address-level identity ([Milionis et al., 2022](#)). Liquid-staking tokens remain separate because they also convey staking returns. Stable assets target a fiat value, imported assets are wrapped representations of off-chain or cross-chain value, and the remaining assets enter the residual category; the asset’s economic exposure determines its category. Stable assets themselves encompass economically different designs. Reserve-backed stablecoins rely on backing and redemption, making confidence in those arrangements relevant to acceptance and run risk ([Gorton and Zhang, 2023](#)); fiat-backed, crypto-collateralized, and algorithmic stablecoins also differ in their solvency and peg-restoration mechanisms ([Catalini et al., 2022](#); [Lyons and Viswanath-Natraj, 2023](#)). Section 3.3 therefore examines the two largest reserve-backed stablecoins separately.

2.3 The route panel

The principal route panel covers Curve, Uniswap v2, v3 and v4, Balancer, SushiSwap v2 and v3, and Fluid on Ethereum. Uniswap v1 enters the separate protocol-architecture sample in Section 5.5. Its retained indexed records identify exchange contracts but omit their underlying token addresses, and the available files contain no complete exchange-to-token crosswalk. Shared transaction hashes and matching ETH legs still recover forced token-to-token routes, while their endpoint-pair identities remain unavailable for the principal panel. Table 1 summarises the principal panel’s coverage. The sample spans 2,332 dates from February 11, 2020 through June 30, 2026. These deployments include several major exchanges but remain a subset of Ethereum trading. Protocols enter after their launch dates, so the exchanges and routes available to traders expand over the sample. We

keep 55 early dates with no observed swaps as zero-activity dates for the covered deployments.

We begin with 472,254,909 pool-level swap records. Indexed blockchain data supply events for seven exchange series, Dune’s `dex.trades` table supplies Fluid records, and direct Ethereum queries supply block headers, ordered logs and receipts, token precision, factory registries, and state checks.⁴ After harmonising trade direction, event identity, and token units, the panel contains 471,616,269 legs. We order those legs within each transaction and apply the flow rule above to distinguish single-pool trades, split orders, sequential routes, and unrelated calls. Appendix A gives the reconstruction algorithm and reports the historical token units that we cannot verify.

Every reconstructed token sequence enters the count measure. Dollar-weighted results require source, intermediary, and destination values to agree within 20%. This is an eligibility test for a coherent dollar notional, not a control for route characteristics: covariates cannot repair a route whose token values do not reconcile. Keeping the count sample broader preserves observed intermediary choice when valuation is noisy; each value exhibit reports the stricter sample’s coverage.

Transactions and contract-state changes are publicly observable (Schär, 2021). Reconstructing a route still requires matching pool events, venues, and token units. This procedure identifies the executed route and its calling contract, but an executor address does not establish which frontend or software formed the quote. Existing router labels likewise cover only part of decentralised-exchange activity (Xi and Moallemi, 2026). On a 2024 snapshot of 74,323 swaps, 241 distinct senders appear, and a registry of the ten largest accounts for 67.7% of those swaps. By a 2025 snapshot, 397 senders appear across 130,282 swaps and the same registry covers 11.8%. We therefore report route-level results without attributing them to a particular routing service.

Table 1: The route panel

	Route panel
Dates in the full route panel	2,332
Missing source-days	0
Raw swap rows	472,254,909
Usable directed atomic trades (legs)	471,616,269
Routes per day, median of 79 sampled days	161,152
Multi-leg routes per day, median of 79 sampled days	31,253
Multi-leg share of routes, median of 79 sampled days	19.7%
Round trips, share of multi-leg routes by count	12.7%
Round trips, share of multi-leg routes by value	21.7%

Note: The full eight-deployment panel contains UTC dates from 11 February 2020 through 30 June 2026. Raw swaps and usable legs use that full panel. The separate Uniswap v1 architecture sample lies outside these totals because its retained events do not recover endpoint-pair token identities uniformly. Daily route statistics and round-trip shares are validation quantities computed on 79 dates sampled across 1 June 2020 through 27 April 2026; this computational sample is not the support of the main dominance estimates. A multi-leg route traverses more than one pool; a round trip ends in the token with which it began. Round-trip shares are formed among multi-leg routes.

⁴See [The Graph’s indexing documentation](#), [Dune’s curated-data documentation](#), and the official [Ethereum node-interface documentation](#).

Table 2: Stablecoin share among native and stable intermediaries

Stablecoin share among native and stable intermediaries	2024	2026	Change [pp] (s.e.)
Route count	16.9%	42.3%	+25.4 (1.05)
Dollar-weighted routes (20% agreement)	32.7%	76.5%	+43.9 (2.02)

Note: The denominator contains exact two-leg routes whose sole intermediary is native or stable; the reported share aggregates all stablecoins in that denominator. Daily shares receive equal weight over the days of the calendar year observed in both years. Dollar-weighted estimates require source, intermediary, and destination values to agree within 20%. Standard errors use a Newey–West covariance with a 30-day Bartlett bandwidth.

3 Rotation in the vehicle role

3.1 A route-level transition


Stablecoins carry a much larger share of vehicle activity in 2026 than in 2024. Comparing the same days of the calendar year, their share among native and stable intermediaries rises from 16.9% to 42.3% by route count and from 32.7% to 76.5% by routed value. The paired-day increases are +25.4 pp and +43.9 pp, respectively. The count series gives equal weight to each exact two-leg route. The dollar-weighted series uses routes for which the values at the source, intermediary, and destination agree within 20%. Table 2 reports both estimates and their uncertainty. Each figure averages over the ultimate pairs active in its year, and those sets differ between the two years. Section 3.2 shows that activity reallocation across ultimate pairs carries the aggregate increase, while positive and negative intermediary changes offset within continuing ultimate pairs.

Figures 1 and 3 extend the comparison to routes of any length and all intermediary types. Each intermediary used on a longer route contributes one observation. In 2026, the 20% value-agreement sample covers 71.6% of native intermediary value and 55.1% of stablecoin intermediary value. The annual path reverses twice. Native assets regain share in 2023 and 2024, stablecoins recover the dollar-weighted lead in 2025, and their lead by route count is shorter lived.

The denominator so far treats a stablecoin as a stablecoin. What backs one moves over the sample, because the collateral an issuer may hold is itself revisable, so each regime label carries a date. Of the 21 stablecoins observed as intermediaries here, 2 change regime within the sample and 0 change between the two years compared. Cutting the same share on the same denominator, calendar, and routing strata by the intermediary’s dated regime therefore holds each regime’s membership fixed across the comparison, and the cut divides the class the way the stablecoin literature does, by what stands behind the peg.

Fiat-reserve tokens move +24.50 pp of the +25.42 pp count rotation, with a standard error of 1.08 pp, and +44.09 pp of the +43.87 pp dollar-weighted rotation, with a standard error of 2.07 pp. The regimes divide one denominator, so their changes add to the class’s, and the residue the fiat-reserve term leaves is small and mixed in sign. Synthetic tokens multiply their share of the class’s routes from 0.3% to 1.4% while contributing +0.55 pp of the count rotation and +1.66 pp of the dollar-weighted one. Tokens that mix real-world assets into on-chain collateral fall −1.73

Figure 1: Intermediary activity rotates between native assets and stablecoins



Note: The sample contains complete routes that do not return to their starting asset. Each route contributes once for every intermediary it uses, so routes with more than two legs can contribute several observations. Shares are formed across all intermediary types; Table 2 uses the narrower native-or-stable denominator. The 2026 observation covers January through June.

pp by value, from 6.8% to 1.4% of the class’s dollars. The rotation is therefore a movement into dollar tokens backed by reserves; the class’s fastest-growing regime moves the aggregate by under two points on either measure. The cut is measured on the released endpoint-vehicle panel; against the daily type panel the pooled rotation differs by 0.001 pp by count and 0.000 pp by value, so the regime terms sit beside the headline in the same units.

3.2 Composition of ultimate pairs

A vehicle’s aggregate share can rise in three ways. It can be chosen more often within an ultimate (source–destination) pair traded in both years; routed activity can move toward continuing ultimate pairs that already use it; or ultimate pairs traded in only one of the two years can route through it. Table 6 shows that the first way accounts for little of the stablecoin increase and that the other two carry it.

Table 2 averages daily shares, whereas Panels A through C of Table 6 pool route count or value across the matched January–June windows. The different weighting explains why the headline changes of +25.4 pp and +43.9 pp differ slightly from the decomposition totals of +25.7 pp and +42.8 pp.

Those three ways form an accounting identity once the third is split in two, because ultimate pairs traded in only one year both route their activity differently and carry a changing share of it. The rest of this section reads the resulting four terms one at a time. Fix one of the measures in Equation (2) and let p index ordered source–destination ultimate pairs. An ultimate pair is continuing when it carries native-or-stable routes in both comparison years and year-specific when it carries them in only one. Let W_y be the share of year- y native-plus-stable mass carried by

continuing ultimate pairs, $\omega_{p,y}$ the weight of continuing ultimate pair p within that mass, and $s_{p,y}$ its stablecoin share, so continuing ultimate pairs route $S_{\text{com},y} = \sum_p \omega_{p,y} s_{p,y}$ of their activity through a stablecoin and year-specific ultimate pairs route $S_{\text{exc},y}$ of theirs. The aggregate share is $S_y = W_y S_{\text{com},y} + (1 - W_y) S_{\text{exc},y}$. Writing $\Delta x = x_{2026} - x_{2024}$ and $\bar{x} = (x_{2024} + x_{2026})/2$,

$$\Delta S = \underbrace{\bar{W} \sum_p \bar{\omega}_p \Delta s_p}_{\text{within continuing ultimate pairs}} + \underbrace{\bar{W} \sum_p \bar{s}_p \Delta \omega_p}_{\text{reweighting}} + \underbrace{(\bar{S}_{\text{com}} - \bar{S}_{\text{exc}}) \Delta W}_{\text{support mass}} + \underbrace{(1 - \bar{W}) \Delta S_{\text{exc}}}_{\text{year-specific ultimate pairs}}. \quad (4)$$

Midpoint weights allocate each product's interaction evenly across its two factors, yielding the four-term identity. Panels B and C of Table 6 report the four terms for route count and routed value, respectively. Within continuing ultimate pairs, the net stablecoin-share change is -0.1 pp by route count and 0.0 pp by value. These small net terms mask offsetting movements; the indented rows split each term into its positive and negative pair contributions. Reweighting across continuing ultimate pairs contributes $+8.6$ pp and $+26.2$ pp; the migration of mass between the continuing and year-specific populations, -0.5 pp and -2.5 pp; and the ultimate pairs traded in only one year, $+17.8$ pp and $+19.2$ pp. The two composition margins carry the rotation between them. Somogyi (2026) decomposes dollar dominance in the other direction, separating the volume of a given atomic dollar pair into the part that exists to route two other currencies and the part that does not. That split runs inside an ultimate pair; the one that carries the rotation measured here runs across them.

Panel A factors the same route-count total a second way. Equation (4) takes an ultimate pair's vehicle-routed activity as given and asks only which vehicle carried it; Panel A asks first how much the ultimate pair traded, then how much of that trading passed through any intermediary at all, and only then which intermediary it was. Let $M_{p,y}$ be the routes observed in ultimate pair p in year y , let $I_{p,y}$ be the fraction of them that passed through a native or stablecoin intermediary, and let $s_{p,y}$ be the stablecoin share of that intermediated activity, the same object the identity above assigns to a continuing ultimate pair. For any set of ultimate pairs $\mathcal{Q}$,

$$F_y(\mathcal{Q}) = \frac{\sum_{p \in \mathcal{Q}} M_{p,y} I_{p,y} s_{p,y}}{\sum_{p \in \mathcal{Q}} M_{p,y} I_{p,y}} \quad (5)$$

is then the stablecoin share those ultimate pairs deliver in year y . Write $\mathcal{N}$ for every ordered ultimate pair, $\mathcal{N}_M \subseteq \mathcal{N}$ for the ultimate pairs that traded in both years, and $\mathcal{N}_V \subseteq \mathcal{N}_M$ for the ultimate pairs that carried an intermediary in both years, so that $S_y = F_y(\mathcal{N})$ and $\mathcal{N}_V$ is the continuing population of Equation (4). Equation (5) supplies the common stable-share functional for these nested populations. Bridging the aggregate to the narrower population twice and allocating what is left inside $\mathcal{N}_V$ gives

$$\Delta S = \underbrace{\Delta[F(\mathcal{N}) - F(\mathcal{N}_M)]}_{\text{ultimate pairs start or stop trading}} + \underbrace{\Delta[F(\mathcal{N}_M) - F(\mathcal{N}_V)]}_{\text{vehicle role starts or stops}} + \underbrace{\phi_M + \phi_I + \phi_s}_{\text{movement inside } \mathcal{N}_V}, \quad (6)$$

where ϕ_M , ϕ_I , and ϕ_s give each factor's average marginal contribution to $F(\mathcal{N}_V)$ over the six orders

in which ultimate-pair activity, vehicle incidence, and stablecoin share can be carried from their 2024 to their 2026 values. Averaging over orders is what lets the three exhaust $\Delta F(\mathcal{N}_V)$ without charging an interaction to whichever factor happens to move first. Read in this order, the classical vehicle-currency question of which intermediary minimises the cost of exchanging a given ultimate pair of assets (Krugman, 1980) is the last of the five terms. The four before it record where trading happens; only the fifth records which intermediary a trade selects.

Ultimate pairs entering or leaving the sample contribute +9.8 pp; established ultimate pairs that gain or lose a vehicle route, -0.4 pp; ultimate-pair activity shifting across the ultimate pairs that carry a vehicle in both years, +7.9 pp; and changes in how often those ultimate pairs use a vehicle at all, +7.0 pp. Direct stable-for-native substitution among them contributes +1.3 pp. The two factorisations agree on which ultimate pairs continue and on nothing else below that. The three allocated terms sum to the change in $\mathcal{N}_V$'s own stablecoin share, while the identity's within-ultimate pair and reweighting terms sum to that same change multiplied by the share of choice mass $\mathcal{N}_V$ carries, averaged over the two years. That relation holds exactly, and it is why a component of one panel is never the same object as the similarly named component of the other. Equation (6) also cuts finer. It splits the ultimate pairs traded in only one year into ultimate pairs that start or stop trading altogether and established ultimate pairs that keep trading and stop or start using a vehicle, and it splits the reweighting of continuing ultimate pairs into how much those ultimate pairs traded and how often they used an intermediary at all. The distance this opens is small where it matters most: the within-ultimate pair term is +1.3 pp under the finer factorisation and -0.1 pp under the identity, and both sit inside the confidence interval of the matched estimate reported in Panel D, whose upper limit is +1.7 pp. Across both count decompositions, positive and negative within-pair changes offset in the aggregate.

The bridge's smallest term is the one that says the most once its population is opened. Ultimate pairs that gain or lose a vehicle route contribute -0.4 pp, which invites the reading that established ultimate pairs barely changed whether they used an intermediary at all. Panel A partitions each year's choice mass into three classes, and the class behind that term is both narrow and closing: ultimate pairs traded in both years whose vehicle role appears in only one carry 3.0% of 2024 choice mass across 4,967 ordered ultimate pairs, and 1.0% of 2026 choice mass across 1,556. Its routed activity passes through a stablecoin more often than either other class does in either year, 47.0% in 2024 and 76.9% in 2026, against 25.0% and 41.2% for the 26,547 ultimate pairs that carry a vehicle in both years and 3.9% and 43.1% for ultimate pairs observed in only one year. The margin on which an established ultimate pair starts or stops using an intermediary at all is therefore narrow, shrinking, and routed through the challenger where it is used. That bounds a third account of the rotation, in which a stablecoin gains aggregate share because ultimate pairs previously observed only on direct routes begin using an intermediary and use a stablecoin when they do. The account covers a population no larger than 3.0% of routed choice mass, far short of a change of +25.7 pp. Membership in the class turns over by construction: an ultimate pair belongs to it only if its vehicle role is present in one year and absent in the other, so no ultimate pair in it carries an intermediary

in both years and none can register a change of intermediary. The distance between the class’s two stablecoin shares records which ultimate pairs occupied it in each year. Appearance in the class and disappearance from it answer to the same forces as the routing decision itself.

Conditional on the same ordered source–destination ultimate pair, calendar date, and single- or cross-exchange route type, stablecoin share changes by +0.2 pp by count and −1.3 pp by value. Panel D of Table 6 reports these estimates and their uncertainty in its all-route count row and supported routed-value row. The corresponding 95% confidence intervals rule out within-ultimate-pair increases larger than +1.7 pp by count and +3.0 pp by value. The matched-ultimate pair estimates are therefore small relative to the aggregate rotation.

Panel D estimates the 2024–2026 difference after absorbing variation specific to an ordered source–destination ultimate pair, calendar date within the year, and observed single- or cross-exchange route type. Let c index each such combination and let $s_{c,y}^{(m)}$ denote its stablecoin share under measure m . We estimate

$$s_{c,y}^{(m)} = \alpha_c^{(m)} + \beta^{(m)} \mathbf{1}\{y = 2026\} + \varepsilon_{c,y}^{(m)}, \quad y \in \{2024, 2026\}. \quad (7)$$

An observation is one ultimate pair–calendar-date–route-type combination in one year. Weighted least squares uses the number or dollar value of native-plus-stable routes as weights; 2024 is the omitted year. Standard errors are clustered by ordered ultimate pair and calendar date. The comparison remains descriptive because the ultimate pairs traded and the use of one or several exchanges can respond to the same forces as intermediary choice.

Conditioning this finely has a price, and the price is coverage. Of the 762,737 ordered ultimate pairs that carry a native-or-stable route in either year, 26,547 carry one in both and make up the continuing block of Equation (4). Requiring the same calendar date and the same route scope as well leaves 5,432 ultimate pairs by count and 5,278 by value. Those ultimate pairs carry 14.0% of 2024 routed activity and 24.3% of 2026, but 41.7% and 39.3% of the dollars moved. Measured inside the continuing block alone, the same four figures are 25.2% and 50.1% by count, 50.8% and 62.9% by value. The matched estimate therefore covers a minority of route counts and a much larger part of the value, because dollars concentrate in the ultimate pairs and days observed on both sides of the comparison. [Amiti et al. \(2022\)](#) read the same gap between a currency’s share of transactions and its share of traded value as evidence about which trades are the small ones; here it says which trades the matched comparison retains. The identity supplies the complement, since its within-ultimate pair term is taken over the whole continuing block under no date or scope condition: −0.1 pp by count and 0.0 pp by value, beside the matched +0.2 pp and −1.3 pp. Each matched interval contains the identity term of its own weighting. The narrow conditioned comparison and the wide accounting term put net within-ultimate pair substitution in the same place, so the offset is not confined to the matched sample.

An entry cut gives the same reading from the first observation of an ultimate pair instead of from a comparison of ultimate pairs active in both years. Among ordered ultimate pairs first observed in January through June 2026, 25.1% of first-window route mass is stablecoin-intermediated; the

analogous 2024 entrant share is 5.0%. The large-notional entry margin is sharper: among entry-day ultimate pair-scope rows with at least \$100,000 of within-20-percent native-or-stable supported value, the stable share rises from 48.4% to 97.8% over 16,686 2026 routes. The increase appears within both venue scopes: single-venue entrant scopes rise from 4.5% to 19.1% and cross-venue scopes from 6.8% to 37.4%. The first observed vehicle also persists. For 2026 entrants with a complete 30-day follow-up window, stable-born ultimate pairs send 97.5% of later native-plus-stable route mass through stablecoins, while native-born ultimate pairs send 0.5%; the gap is +97.0 pp. At 120 days the corresponding gap remains +98.8 pp. The controlled estimates preserve the contrast. For each entry horizon $h \in \{30, 120\}$, the controlled specification estimates

$$\begin{aligned}
S_{p,t+h} = & \alpha_h \\
& + \beta_h S_{p,t}^{\text{entry}} \\
& + \gamma_h \mathbf{1}\{S_{p,t}^{\text{entry}} > 1/2\} \\
& + X'_{p,t} \delta_h \\
& + \varepsilon_{p,t+h},
\end{aligned} \tag{8}$$

where $S_{p,t+h}$ is the follow-up stable-vehicle share for non-WETH entrant ultimate pair p , $S_{p,t}^{\text{entry}}$ is the first-window stable share, and X contains entry size, cohort year, endpoint class, direct-route share, and route-complexity share. Among those non-WETH entrants, a 1 pp higher entry-day stable share predicts +0.86 pp more stable follow-up share at 30 days and +0.74 pp at 120 days; stable-dominant entry predicts +57.3 pp higher probability of stable-majority follow-up at 120 days. The supported split is mainly the no-direct-route case: the 120-day panel has 512,020 no-direct-route entrant rows and 181 direct-route-present rows. Inside the no-direct margin, a 1 pp higher entry stable share predicts +0.89 pp more stable follow-up share. The value-supported analogue is similar: a 1 pp higher entry stable-value share predicts +0.90 pp more follow-up stable-value share at 30 days and +0.78 pp at 120 days, in value-weighted non-WETH entrant rows. Stable birth is also related to continued ultimate-pair activity: among 512,201 non-WETH entrants, stable-present birth predicts +9.2% more non-entry-day primary routes over 120 days and +5.9 pp higher probability of trading again. Active-day histories show the same persistence. Among 2026 stable-dominant entrants that trade again, 92.2% never register a non-stable-majority active day over the first month, and 95.1% never do so in the complete 120-day sample. The result extends beyond WETH endpoints: WETH-endpoint entrants carry 18.9% of 2026 entrant route mass, and among entrants without WETH at an endpoint stable share still rises from 0.9% to 7.7%. The secure-volume mechanism in [Krugman \(1980\)](#) is visible within that non-WETH margin. In 2026, entrants with a stablecoin already at an endpoint route 17.5% of their first-window mass through stablecoins, versus 5.2% for other non-WETH endpoints, and the stable-endpoint gap expands by +10.3 pp from 2024 to 2026. After entry size, endpoint class, direct-route share, and route-complexity controls, the 2026 stable-endpoint interaction is +2.4 pp. Two issuers concentrate the stable-entrant margin: in 2026 USDC carries 73.8% and USDT 25.8% of stable-entry route mass, 99.6% together. Identity

persists within each issuer: USDC-born stable entrants keep 90.5% of stable follow-up routes in USDC over 30 days and 88.3% over 120, while the corresponding USDT shares are 77.8% and 87.7%. The controlled vehicle-level analogue uses 6,087 non-WETH stable-entry vehicle rows with asset fixed effects. Conditional on positive stable follow-up mass, a 1 pp higher entry share for a named stablecoin predicts +0.98 pp more own-vehicle follow-up share at 30 days and +0.97 pp at 120 days. In a weighted entrant regression that drops WETH-endpoint ultimate pairs, the 2026 coefficient remains +4.3 pp after endpoint class, entry size, direct-route, and route-complexity controls; the additional 2026 stable-endpoint term is +5.7 pp. A route-architecture analysis points in the same direction: the 2026 interaction with direct-route share is +50.2 pp, and the 2026 interaction with complex-route share is +35.5 pp. These entry, endpoint, architecture, and identity estimates are the entry analogue of the selection margin in [Amiti et al. \(2022\)](#): composition can move an aggregate while vehicle identities inside continuing units remain persistent. The estimates describe persistence after entry; causal effects of the first route and displacement of an active direct atomic-pair venue remain outside the design. New ultimate pairs, especially those that require an intermediary, tend to arrive attached to a vehicle regime that persists over the first four months.

The same identity persistence appears in a stress episode. During the March 2023 USDC/SVB window, stable-vehicle route activity rose by 73.6% to 8,931 routes per day, but USDC’s share inside the stable-vehicle class was 67.9%, against 68.0% in the prior window, a change of -0.2 pp. USDT’s share rose only afterward, to 25.0% in the post-window. The capital side diverged from the route shock: inside the five-vehicle set, stablecoins’ route share rose by +11.4 pp during the window while their deposited-capital share moved -2.0 pp, and after the window WETH’s capital share was +3.1 pp higher than before. Stablecoin pegs depend on arbitrage and redemption arrangements ([Lyons and Viswanath-Natraj, 2023](#); [Anadu et al., 2025](#)), so this peg shock remains observational. Even so, the episode is informative about identity under stress: USDC continued to carry the vehicle role while its cash-like promise was under pressure, and LP capital shifted toward the native reserve stock during the temporary stable-route surge.

The entry result extends to repeatedly traded ultimate pairs. On the same matched January–June calendar, stablecoin share among ultimate pair-days with at most five realised routes rises from 6.9% to 20.9%, with a change of +14.0 pp and standard error 1.1 pp. Among ultimate pair-days with more than 100 realised routes, the stable share is already 64.8% in 2024 and rises to 73.8% in 2026. Active ultimate pair-days carry +52.9 pp more stable share than thin ultimate pair-days in 2026. The same check at ultimate-pair entry is sharper. Entrants with more than 100 first-day routes carry 20.2% of 2026 entrant route mass and route 91.4% through stablecoins, up by +32.6 pp from 2024. Thin entrants in 2026 route only 6.2% through stablecoins. Because realised route activity is an outcome, this size split is descriptive; it shows that the composition margin lives only in isolated low-activity ultimate pairs.

In [Table 3](#), the Entry persistence, No-direct-route split, Value-weighted entry, and Named-stable identity rows report the persistence estimates; the Secure-volume endpoint and Route architecture rows report the formation margins.

Table 3: Ultimate-pair entry and later vehicle use

Specification	Outcome	Regressor	Coefficient	Obs. / clusters
Entry persistence	Stable follow-up share ($S_{p,t+h}$), 30 days	Entry stable share ($S_{p,t}^{\text{entry}}$)	+0.86*** (0.06)	599,870 / 333
Entry persistence	Stable follow-up share ($S_{p,t+h}$), 120 days	Entry stable share ($S_{p,t}^{\text{entry}}$)	+0.74*** (0.13)	512,201 / 243
Entry persistence	Stable-majority follow-up, 120 days	Stable-majority entry ($\mathbf{1}\{S_{p,t}^{\text{entry}} > 1/2\}$)	+57.3*** (14.8)	512,201 / 243
No-direct-route split	Stable follow-up share ($S_{p,t+h}$), 120 days	Entry stable share ($S_{p,t}^{\text{entry}}$), no direct route	+0.89*** (0.13)	512,020 / 243
Value-weighted entry	Stable value share, 120 days	Entry stable value share	+0.78*** (0.14)	427,625 / 243
Named-stable identity	Own-stable follow-up share, 120 days	Entry named-stable share	+0.97*** (0.02)	6,087 / 243
Secure-volume endpoint	Entry stable share ($S_{p,t}^{\text{entry}}$)	Stable endpoint $\times$ 2026	+5.7*** (1.3)	615,333 / 363
Route architecture	Entry stable share ($S_{p,t}^{\text{entry}}$)	Direct-route share ($D_{p,t}$) $\times$ 2026	+0.50*** (0.14)	615,333 / 363
Route architecture	Entry stable share ($S_{p,t}^{\text{entry}}$)	Complex-route share ($C_{p,t}$) $\times$ 2026	+0.36*** (0.12)	615,333 / 363

Note: Each cell reports a regression coefficient with entry-date-clustered CR1 standard errors in parentheses. In the rows governed by Equation (8), $S_{p,t+h}$ is stablecoin-intermediated follow-up route mass divided by native-plus-stable follow-up route mass through horizon h , and $S_{p,t}^{\text{entry}}$ is the same share in the first observed window. Share outcomes and share regressors are measured in pp, so a share-on-share coefficient is the pp change in the outcome for a 1 pp increase in the regressor; stable-majority entry and stable-endpoint coefficients are pp differences for binary regressors. Direct-route and complex-route rows are 2026 interactions. The unit is a non-WETH entrant ultimate pair-scope row except for the named-stable identity row, which is a stablecoin entry row with asset fixed effects. Rows are route-weighted except the stable-value row, which uses routed value whose source, intermediary, and destination values agree within 20%. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively.

The transition panel next asks which initial activity and route-architecture states are associated with the 2024–2026 rotation. For measure m , ordered ultimate pair p , calendar month-day d , and realised route scope s , the common design is

$$\Delta S_{pds}^{(m)} = X'_{pds,2024} \beta^{(m)} + \Delta Z'_{pds} \gamma^{(m)} + \mu_d^{(m)} + \varepsilon_{pds}^{(m)}, \quad (9)$$

In Equation (9), $m \in \{N, V\}$ indexes route-count and supported routed-value shares. For each transition cell, $N_{pds,y}$ is its number of realised routes, $D_{pds,y}$ is the share that routes directly, $C_{pds,y}$ is the share using multiple intermediaries, split or joined legs, or nonsequential atomic trades, $P_{pds,y}$ is native-plus-stable primary-choice route count divided by total route count, $A_{p,y}$ is ultimate-pair age in days, and V_{pds} indicates a cross-exchange route scope. Thus $X_{pds,2024}$ contains $\log(1 + N_{pds,2024})$, $D_{pds,2024}$, $C_{pds,2024}$, $P_{pds,2024}$, $\log(1 + A_{p,2024})$, and V_{pds} , while ΔZ_{pds} contains the 2024–2026 changes in log route count, direct-route share, and complex-route share. Following the published-panel convention in Somogyi (2026), each panel of Table 4 keeps the dependent variable, population, fixed effects, weighting, and covariance design fixed while columns add initial activity, route architecture, their joint specification, and changes in both. Month-day effects absorb the common calendar; standard errors cluster by ordered ultimate pair and month-day; harmonic endpoint route mass weights the two endpoint years symmetrically. These descriptive determinants summarise realised associations under the stated fixed effects.

A second analysis examines stablecoin entry from native-only baselines. The benchmark is

Table 4: Determinants of stablecoin vehicle-share growth

	(1) Market state	(2) Route architecture	(3) Joint baseline	(4) + dynamics
<i>Panel A: $\Delta S_{pds}^{(N)}$, route-count stable share [pp]</i>				
$\log(1 + N_{pds,2024})$	-0.571** (0.236)		+0.186 (0.339)	-0.006 (0.344)
$D_{pds,2024}$		-3.589 (5.542)	-3.985 (5.762)	-1.119 (6.135)
$C_{pds,2024}$		-2.391 (7.202)	-1.859 (6.901)	-1.025 (7.914)
$P_{pds,2024}$		+0.940 (4.632)	+1.568 (4.043)	+2.367 (4.467)
$\log(1 + A_{p,2024})$	-0.296 (0.291)		+0.116 (0.230)	-0.081 (0.229)
V_{pds}	-1.115 (0.693)	-1.339* (0.748)	-1.367* (0.761)	-1.299* (0.763)
$\Delta \log(1 + N_{pds})$				-0.007 (0.366)
ΔD_{pds}				+5.301** (2.651)
ΔC_{pds}				+0.024 (1.857)
Within R^2	0.006	0.010	0.010	0.013
Observations	94,260	94,260	94,260	94,260
Ordered ultimate-pair clusters	5,432	5,432	5,432	5,432
Month-day clusters	181	181	181	181
Month-day fixed effects	Yes	Yes	Yes	Yes
Harmonic route-mass weights	Yes	Yes	Yes	Yes
<i>Panel B: $\Delta S_{pds}^{(V)}$, routed-value stable share [pp]</i>				
$\log(1 + N_{pds,2024})$	+0.129 (0.473)		+1.270 (1.127)	+0.327 (0.790)
$D_{pds,2024}$		-20.123 (19.987)	-13.388 (15.754)	-0.760 (16.215)
$C_{pds,2024}$		-9.485 (22.402)	+2.783 (19.155)	+5.596 (19.827)
$P_{pds,2024}$		-12.979 (18.315)	-2.899 (14.636)	-2.952 (14.776)
$\log(1 + A_{p,2024})$	-2.471 (1.604)		-1.761 (1.192)	-2.155* (1.260)
V_{pds}	-5.016 (3.503)	-5.084 (3.604)	-5.090 (3.479)	-4.456 (3.077)
$\Delta \log(1 + N_{pds})$				-2.122* (1.205)
ΔD_{pds}				+23.220** (11.559)
ΔC_{pds}				+9.839 (6.061)
Within R^2	0.019	0.021	0.030	0.047
Observations	91,417	91,417	91,417	91,417
Ordered ultimate-pair clusters	5,278	5,278	5,278	5,278
Month-day clusters	181	181	181	181
Month-day fixed effects	Yes	Yes	Yes	Yes
Harmonic route-mass weights	Yes	Yes	Yes	Yes

Note: The unit is an ordered ultimate-pair-calendar-month-day-route-scope transition cell observed in both 2024 and 2026. The dependent variables $\Delta S_{pds}^{(N)}$ and $\Delta S_{pds}^{(V)}$ are changes in route-count and supported routed-value stablecoin shares, respectively; supported value requires source, intermediary, and destination dollar values to agree within 20%. The regressors are defined immediately after Equation (9); Δ denotes the 2026 value minus the 2024 value. Coefficients and clustered standard errors in parentheses are in pp. All columns absorb month-day effects, use harmonic native-plus-stable route-mass weights, and cluster by ordered ultimate pair and month-day. Columns (1) and (2) enter the initial-activity and route-architecture blocks separately; column (3) combines their baseline variables; column (4) adds changes in ultimate-pair activity and route architecture. Within R^2 is computed after the fixed-effect transformation. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively. The estimates are descriptive associations.

regime persistence inside continuing ultimate pairs: among single-venue transition cells, 2024 stable-majority regimes remain stable-majority in 94.5% of 2026 observations, and native-majority regimes remain native-majority in 95.5%. In the count-share transition panel, a one-log-point larger baseline route count is associated with a -2.7 pp lower turn-on probability, and a one-standard-deviation increase gives -3.7 pp; stable-majority switches move the same way, -1.1 pp. Baseline direct-route and complex-route shares enter positively, $+27.1$ pp and $+20.3$ pp. The direct-route association is stronger among lower-activity ultimate pairs: the direct-share by one-log-point-thinner interaction is $+3.1$ pp. The first five rows of Table 5 report these turn-on, thin/direct, and leader-switch estimates.

A daily native-only hazard gives the timing version of the same margin. Among 1,690,478 active ultimate pair-days with no stable vehicle at the origin date, the probability that a stable vehicle appears over the next 30 days is 24.5% in the highest ultimate-pair-activity bin and 2.0% in the lowest, a gap of $+22.5$ pp. The oldest native-only ultimate pairs show 41.4% turn-on against 1.6% for the youngest, a gap of $+39.9$ pp. With month-day fixed effects, ordered ultimate-pair and date clustering, and ultimate-pair route-count weights, one log point of origin route count predicts $+3.5$ pp higher 30-day turn-on probability; ultimate-pair age predicts $+1.6$ pp. These coefficients are the two Rolling hazard rows of Table 5. Stable vehicles appear after an ultimate pair has repeatedly produced route demand; within mixed same-day sets, native vehicles retain the larger share.

A mixed risk-set check makes the last point directly. Among observed ultimate pair-day-scope sets in which native and stable vehicles both appear and at least five routes trade, stable vehicles carry 28.3% of 2026 routes, up from 24.3% in 2024. With ultimate pair-day-scope fixed effects, however, a stable row carries -51.9 pp less route share than a native row; the 2026 interaction is $+3.7$ pp. Vehicle network reach explains more of the within-set allocation. Adding the vehicle’s same-day ultimate pair-scope count outside the current risk set, one log point of reach predicts $+9.4$ pp more route share, while the stable-row coefficient remains -21.5 pp and the 2026 stable interaction is $+2.3$ pp. Splitting stable rows by issuer sharpens the boundary: after the same risk-set effects and reach control, the generic stable-2026 term is -40.9 pp, while the USDC and USDT 2026 issuer interactions are $+53.0$ pp and $+88.2$ pp. Prior-30-day reach produces the same pattern: replacing same-day reach with the vehicle’s ultimate pair-scope reach over the prior 30 calendar days gives $+7.8$ pp, and the stable-row coefficient remains negative at -20.4 pp. These predictive specifications use observational variation. Together they place formation at the edge of incumbent regimes and in ultimate pairs with repeated demand, while pointing away from simple same-risk-set conquest.

The Mixed risk set and Issuer split rows of Table 5 report the reach and issuer interactions.

The joint condition is easiest to read in the estimator’s own unit. An observation is one ultimate pair–calendar-date–route-type cell, and 94,260 cells clear the condition, out of the 1,726,215 that carried a native-or-stable route in 2024 and the 815,483 that carried one in 2026, or 5.5% and 11.6% of the two years’ cells. Those cells carry 14.0% and 24.3% of the two years’ routes, so a matched cell is 2.6 times as busy as the average active cell in 2024 and 2.1 times as busy in 2026. On

Table 5: Stable-vehicle turn-on and network reach

Margin	Outcome	Regressor	Coefficient [pp]	Obs. / clusters
Turn-on	Stable appears	Baseline route count	-2.7*** (0.4)	94,260 / 181
Turn-on	Stable appears	Direct-route share	+27.1*** (6.4)	94,260 / 181
Turn-on	Stable appears	Complex-route share	+20.3*** (5.1)	94,260 / 181
Thin/direct	Stable appears	Direct share $\times$ thinness	+3.1*** (0.9)	94,260 / 181
Leader switch	Stable becomes leader	Baseline route count	-1.1*** (0.2)	94,260 / 181
Rolling hazard	Stable appears within 30d	Origin route count	+3.5*** (0.7)	1,690,478 / 303
Rolling hazard	Stable appears within 30d	Ultimate-pair age	+1.6*** (0.2)	1,690,478 / 303
Mixed risk set	Vehicle route share	Same-day reach	+9.4*** (2.9)	27,852 / 363
Mixed risk set	Vehicle route share	Prior-30d reach	+7.8*** (2.5)	27,852 / 363
Issuer split	Vehicle route share	USDC $\times$ 2026	+53.0** (22.5)	27,852 / 363
Issuer split	Vehicle route share	USDT $\times$ 2026	+88.2*** (24.2)	27,852 / 363

Note: Each cell reports a predictive coefficient in pp, with clustered standard errors in parentheses. Turn-on is an indicator for zero stablecoin route share in 2024 and positive stablecoin route share in 2026; leader switch is an indicator for stablecoin route share crossing one half over the same interval. Rolling hazard indicates that an active native-only ordered ultimate pair gains a stablecoin route within 30 days. Route count, ultimate-pair age, same-day reach, and prior-30-day reach enter as $\log(1+x)$, so their coefficients correspond to one log point. Direct-route and complex-route shares enter as unit shares; issuer interactions are binary. Turn-on and leader-switch rows use month-day effects and harmonic endpoint-denominator weights; rolling-hazard rows have complete look-ahead support and cluster by ordered ultimate pair and date; mixed-risk-set and issuer rows absorb ultimate pair-day-scope risk-set effects, use route-count weights, and cluster by ordered ultimate pair and date. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively. The estimates are predictive associations from observational variation.

the dollar-weighted perimeter, where the value-agreement requirement empties 142,972 cell-years that the count measure keeps, 5.7% and 11.6% of cells are matched, and a matched cell carries 7.4 times the average cell’s dollars in 2024 and 3.4 times in 2026. A cell survives only when the same ultimate pair traded on the same day of the calendar year in both years, and the chance of that rises with how often the ultimate pair trades. The matched estimate is therefore a statement about the subset of ultimate-pair cells that trades routinely, and the dollar multiple says that this part is where value concentrates. It speaks less for ultimate pairs that trade intermittently, whose cells drop out even when the ultimate pair itself trades in both years. That division suits the question: the rotation is a fact about where dollars route, and the matched comparison is measured on the cells that carry two-fifths of them in either year.

Individual ultimate pairs make the accounting concrete. Among ultimate pairs whose two assets carry a recognised ticker, the largest contribution to the reweighting term comes from $\text{USDT} \rightarrow \text{WETH}$, which is stablecoin-intermediated in both years and grows from 7,447 to 54,112 routes, or from 0.36% to 4.52% of routed activity, and adds +2.17 pp on its own. The largest such contribution among ultimate pairs absent from 2024 comes from $\text{USDS} \rightarrow \text{WETH}$, which supplies 3,390 stablecoin-intermediated routes in 2026 and adds +0.13 pp. The largest such contribution within a continuing ultimate pair, from $\text{USDC} \rightarrow \text{WBTC}$, adds +0.05 pp as its stablecoin share moves from 10.5% to 35.9%; the within-ultimate pair term summed over all ultimate pairs is -0.1 pp. Because most newly traded assets have no such ticker, each named ultimate pair understates

Table 6: Decomposition of the stablecoin rotation

Margin or estimate	Estimate [pp]	Obs.
<i>Panel A. Route-count share: market activity and vehicle incidence</i>		
Ultimate pairs entering or leaving the sample	+9.8	
Ultimate pairs gaining or losing a vehicle route	−0.4	
Market activity shifting across continuing ultimate pairs	+7.9	
Change in how often continuing ultimate pairs use a vehicle	+7.0	
Stablecoin share within continuing vehicle-using ultimate pairs	+1.3	
Total route-count change	+25.7	
<i>Panel B. Route-count share: continuing and year-specific ultimate pairs</i>		
Net stablecoin-share change within continuing ultimate pairs	−0.1	
Pairs moving toward stablecoins (1,575)	+1.3	
Pairs moving toward native assets (1,489)	−1.4	
Vehicle activity shifting across continuing ultimate pairs	+8.6	
Weight of continuing versus year-specific ultimate pairs	−0.5	
Ultimate pairs traded in only one year	+17.8	
Total route-count change	+25.7	
<i>Panel C. Dollar-weighted share: continuing and year-specific ultimate pairs</i>		
Net stablecoin-share change within continuing ultimate pairs	0.0	
Pairs moving toward stablecoins (1,511)	+2.3	
Pairs moving toward native assets (1,445)	−2.4	
Vehicle activity shifting across continuing ultimate pairs	+26.2	
Weight of continuing versus year-specific ultimate pairs	−2.5	
Ultimate pairs traded in only one year	+19.2	
Total change in dollar-weighted share	+42.8	
<i>Panel D. Matched ordered ultimate-pair estimates</i>		
All two-leg routes, count share	+0.22 (0.76)	188,520
20% agreement sample, count share	+0.32 (0.75)	182,834
20% agreement sample, dollar-weighted share	−1.35 (2.19)	182,834

Note: Panel A decomposes the route-count share change between 2024 and 2026 over observed ultimate-pair activity, vehicle incidence, and stablecoin share; common-role ultimate pairs use a native or stablecoin intermediary in both years, and a Shapley allocation distributes interactions across the five reported components. Panels B and C report the four terms of Equation (4), which conditions on native-plus-stable choice mass throughout; their indented rows split the net within-pair term into positive and negative pair contributions and are not additional decomposition components. Panel C covers dollar-weighted routes whose source, intermediary, and destination values agree within 20%. Panels A and B decompose the same route-count total by different factorisations, so their components are not comparable term by term. Panel D estimates Equation (7); standard errors use two-way clustering by ordered ultimate pair and calendar date.

its margin, which is why the margin totals accompany it. New listings or changes in pool access can bring an ultimate pair into the sample; trading activity can shift across continuing ultimate pairs; and direct atomic-pair depth or router access can change the use of indirect execution. Through these margins, routed activity reallocates across source–destination ultimate pairs. This reallocation carries most of the stablecoin increase. Net stable-for-native substitution within continuing vehicle-using ultimate pairs contributes far less.

A margin of a given size can come from one ultimate pair or from thousands of ultimate pairs, and the totals do not distinguish the two. It is therefore natural to ask whether a few large ultimate pairs carry the reallocation. The two composition margins answer differently. Ultimate pairs traded only in 2026 contribute through 18,446 ordered ultimate pairs, the largest of which supplies 3.2% of that margin, and 1,386 of them are needed for nine-tenths of it. Reweighting among continuing ultimate pairs is top-heavy instead: 7 ultimate pairs supply half of its +17.2 pp gross gain, and the margin nets to +8.6 pp only because 2,817 continuing ultimate pairs lose activity weight, −8.6 pp between them. Dollar weighting sharpens the same contrast, with 5 ultimate pairs behind half of a +38.0 pp gross gain. Within continuing ultimate pairs, intermediary choice moves in both directions: 1,575 of them raise their stablecoin share and 1,489 lower it. The ultimate pairs that rise contribute +1.3 pp by count and +2.3 pp by value; the ultimate pairs that fall give back −1.4 pp and −2.4 pp. Positive and negative changes in intermediary choice offset across continuing ultimate pairs, leaving their net contribution small beside the activity weight that moves across ultimate pairs. The rotation therefore travels along two different edges of the trading network: the network extends into many new and mostly small ultimate pairs, while activity weight concentrates in a few ultimate pairs that were already large. This reorganisation of where trading happens carries the aggregate change; the growth of a single ultimate pair does not. The two edges also separate what a currency network measures: a matrix recording which ultimate pairs trade at all, as in the historical exchange structure of [Flandreau and Jobst \(2009\)](#), registers the first margin and is silent on the second.

The vehicle role is not contestable everywhere in the ultimate pair space. A route whose source or destination is wrapped ether cannot also use wrapped ether as its intermediary, so within the native-versus-stablecoin comparison such a route is stablecoin-intermediated by definition, in 2024 and in 2026 alike. Since the ultimate pairs that grew fastest are concentrated among exactly those ultimate pairs, a question of interest is whether the two composition margins record anything more than the growth of ultimate pairs whose vehicle was never in doubt. We split each margin according to whether wrapped ether is an endpoint of the ordered ultimate pair. One contract address determines the split, giving the long tail of newly traded assets the same classification precision as the majors. 1,499 of the 26,547 continuing ultimate pairs, or 5.6%, have such an endpoint, and their share of routed activity rises from 20.9% to 33.0%. By route count they supply +6.3 pp of the +8.6 pp reweighting margin and +13.9 pp of the +21.0 pp new-ultimate pair margin, so most of the count rotation does run through ultimate pairs that had no intermediary to choose. Dollar weighting reverses the reading. The same ultimate pairs supply +11.2 pp of

the +26.2 pp value reweighting margin, or 42.6% of it, and +3.6 pp of the value new-ultimate pair margin, or 16.0%. Once trades are weighted by the value they move, the rotation is carried by ultimate pairs that could have been routed either way and were routed through a stablecoin. The same identity disciplines the within-ultimate pair term: an ultimate pair with a wrapped-ether endpoint contributes exactly zero to it under both weightings, so ultimate pairs fixed at a stablecoin share of one do not mechanically dilute the total. Every unit of it comes from ultimate pairs whose intermediary was a live choice, and their positive and negative changes continue to offset in aggregate.

So far the split has been taken over all routes at once, which does not say where in the venue structure the ultimate pairs without a vehicle choice sit. Taking it separately for routes that stay on one exchange and routes that cross exchanges answers that, and the two weightings answer it differently. Within a single exchange, reweighting among continuing ultimate pairs contributes +4.4 pp, of which wrapped-ether-endpoint ultimate pairs supply 55.3%. Across exchanges the same margin is +11.4 pp, and those ultimate pairs supply 82.9% of it. What survives the split is the part contributed by ultimate pairs that had a choice, and it is the same size in both: +2.0 pp within an exchange and +1.9 pp across exchanges. The entire cross-exchange excess in the count margin is therefore ultimate pairs that were stablecoin-intermediated by definition, and the ultimate pairs that could have gone either way reallocate activity at the same rate wherever the route runs. Weighting by the value moved removes the contrast. The margin is +26.6 pp within an exchange against +21.4 pp across exchanges, and eligible ultimate pairs supply 48.6% and 50.5% of them. In route counts, then, the rotation looks like something that happened to cross-exchange trading, and what it records there is which ultimate pairs became active across exchanges; in value terms it is no more a cross-exchange event than a single-exchange one. [Makarov and Schoar \(2020\)](#) decompose bitcoin price dispersion the same way and find the across-region component both the larger one and, once traded volume is taken into account, the one that carries the economic magnitude. The across-venue excess here does the opposite: it is visible in unweighted route counts and gone from the dollar-weighted margin.

That split was taken of continuing ultimate pairs, whose endpoints are fixed and whose intermediary can change. The larger composition margin comes from ultimate pairs that traded only in 2026, and it is there that a reader should most expect eligibility to be the whole story: an ultimate pair appearing for the first time is disproportionately a newly listed asset whose only quote is a wrapped-ether ultimate pair, which is to say an ultimate pair that arrives with its intermediary already settled. By route count that expectation is met, and the route scope barely moves it. Entry contributes +21.4 pp within a single exchange and +23.8 pp across exchanges, and wrapped-ether-endpoint ultimate pairs supply 65.8% of the first and 67.5% of the second. Dollar weighting separates the two scopes instead of reversing them. The entry margin is +14.1 pp within an exchange and twice that, +28.5 pp, across exchanges, while the part from ultimate pairs that had no vehicle to choose is +3.4 pp and +4.0 pp, almost the same in both. What differs is the part that had a choice: +10.7 pp within an exchange and +24.5 pp across it, so eligible ultimate

pairs account for 24.1% of the margin within an exchange and 14.2% across it. The entry margin is therefore the mirror image of the reweighting margin. There the cross-exchange excess was carried entirely by ultimate pairs whose vehicle was never in question; here it is carried entirely by ultimate pairs that had one and were routed through a stablecoin anyway. Read together, the two say that an ultimate pair already trading across exchanges gains stablecoin share mainly where no other intermediary was available, while the value arriving with ultimate pairs that begin to trade across exchanges comes from ultimate pairs that could have been routed through the native token.

Those splits sort ultimate pairs by whether the vehicle was fixed and by where the route ran. The economically live comparison is ultimate pairs where either vehicle class can carry the route. We divide them once more, on whether a stablecoin stands at an endpoint of the ordered ultimate pair. The stable-endpoint split preserves two feasible intermediary classes: a trade into one dollar token can be carried by wrapped ether or by a different dollar token. The 10,278 continuing ultimate pairs of this class route through a stablecoin in 5.6% of their 2024 activity and 12.8% of their 2026 activity, well below the unit share eligibility imposes, and they hold 68.8% and 60.5% of routed activity in the two years. By route count they supply 26.0% of the reweighting margin and 28.7% of the entry margin, less in both cases than wrapped-ether ultimate pairs supply. Dollar weighting puts them ahead. The same ultimate pairs carry +15.0 pp of the +26.2 pp value reweighting margin, or 57.3% of it, and +16.9 pp of the value entry margin, or 76.4%, both larger than the eligible ultimate pairs' 42.6% and 16.0%. Ultimate pairs holding neither vehicle at an endpoint are the many and the small: 142,831 of them begin trading in 2026 and 13,390 trade in both years, and they supply +1.7 pp and +0.0 pp of the two value margins. Inside the dollar-endpoint ultimate pairs the movement is once again reallocation. Their activity-weighted stablecoin share by value rises from 21.0% to 61.6% while their within-ultimate pair term is -0.0 pp, so the rise is weight moving toward ultimate pairs that already routed through a dollar. Weighted by the value they move, the ultimate pairs that carry the rotation have a dollar already at one end, and the intermediary inside each of them barely moved on net. [Krugman \(1980\)](#) calls the trading a currency draws from payments that never pass through it its secure volume, and holds that a currency short of it cannot take the vehicle role from an incumbent however transaction costs fall. The ultimate pairs carrying this rotation are the ones where the challenger's secure volume sits.

That entry margin is gross. The decomposition subtracts the ultimate pairs that stopped trading, and the netting is integral to the comparison. Averaged across the two years, ultimate pairs observed in only one of them carry 48.0% of route activity and 27.8% of the value moved, and the netted term is that mass multiplied by the gap between how two populations were routed—the ultimate pairs that left and the ultimate pairs that arrived. The two margins have different economic content. The departing ultimate pairs sent 6.8% of their routes and 10.8% of their dollars through a stablecoin; the ultimate pairs that replaced them send 43.8% and 79.8%. Removing the wrapped-ether-endpoint ultimate pairs, whose vehicle was settled before they opened, leaves the departing ultimate pairs at 1.7% by route count and 3.2% by value, and the arriving ones at 21.0% and 76.9%. Across exchanges the arriving ultimate pairs that could have chosen anything route

88.3% of their dollars through a stablecoin. Ultimate-pair entry is the largest margin in the count decomposition. A new population replaces the departing ultimate pairs and routes differently from the day it opens. [Flandreau and Jobst \(2009\)](#) estimate the incumbency advantage of an international currency across a stable roster of trading relationships and find persistence without lock-in; here the roster is what turns over, and the challenger’s gain is concentrated in the ultimate pairs that did not exist when the incumbent’s network was formed.

Two different economies produce that gap, and the two weightings do not agree on which one is at work. Ultimate pairs are of three kinds: those with wrapped ether at an endpoint, which cannot use it to intermediate and route through a stablecoin by construction; those with a stablecoin at an endpoint, which may still be carried either way; and those with some other asset at each end. Evaluated at the two cohorts’ midpoints, the gap separates exactly into the mix of kinds each cohort brings and the rate at which each kind routes. By route count the mix carries 65.9% of the +17.8 pp netted term, and nearly all of that is the first kind: wrapped-ether-paired ultimate pairs are 28.9% of arriving activity against 5.2% of departing activity, worth +11.4 pp by themselves. They enter the margin through their weight alone, because their routing rate is one in both cohorts. The remaining +6.1 pp is a rate difference among ultimate pairs that could have been carried either way. By value the ordering reverses: the mix contributes +3.0 pp and the rates +16.2 pp, or 84.6% of +19.2 pp. Arriving ultimate pairs with a stablecoin at an endpoint send 81.7% of their dollars through a stablecoin where the departing ultimate pairs of the same kind sent 6.3%, and arriving ultimate pairs with some other asset at each end send 48.3% against 0.9%. No ultimate pair appears on both sides of those comparisons, so none of these figures describes an intermediary being changed; the dollars now moving through ultimate pairs that opened after the challenger was established were never routed the old way. In [Dowd and Greenaway \(1993\)](#), willingness to switch depends on the expectation that others will switch too, and the resulting excess inertia applies to agents already holding the incumbent. Their model is silent on the agents who arrive later, and it is those the data separate: inside a continuing ultimate pair the intermediary barely moves, while the ultimate pairs that open once the challenger is available carry it from their first trade.

Replacement on that scale moves activity between the two populations as well as changing how each is routed, and the decomposition prices the two separately. Ultimate pairs observed in both years fell from 55.5% of route activity to 48.6%, a shift of -6.9 pp, and from 82.0% of the value moved to 62.4%, a shift of -19.6 pp; across exchanges the dollar shift is -27.6 pp. Roughly a fifth of the value in the comparison therefore changed which population of ultimate pairs carried it. Against those magnitudes the term itself is -0.5 pp by route count and -2.5 pp by value, the smallest of the four and the only one no other margin overlaps. Large activity shifts produce a small term because the two populations are priced almost alike: averaged over the two years, ultimate pairs alive in both route 33.1% of their routes and 58.2% of their dollars through a stablecoin, against 25.3% and 45.3% for ultimate pairs alive in only one, gaps of +7.8 pp and +12.9 pp. The term is negative in every weighting and every route scope: holding both populations at those midpoint routing rates, the activity that changed population lowers the stablecoin share. Activity reallocation within the

continuing-ultimate-pair block produces the positive rotation because the population receiving the activity routes differently from the one it replaced. [Amiti et al. \(2022\)](#) find invoicing currency to be close to a fixed attribute of an exporter, switching for a small minority of observations over almost four years, with the composition of destinations explaining little of the cross-sectional variation. Intermediary identity is similarly persistent inside a continuing ultimate pair here; the difference is that the population of ultimate pairs turns over fast enough for composition to carry the whole aggregate change.

Composition is measured on a population that turns over, and the same measurement can be taken where it holds still. Fix a set of ultimate pairs and the identity’s four terms collapse to the two that operate inside it. The two readings of the rotation then predict opposite outcomes. Under turnover, holding the set fixed removes the rotation. Under reallocation among ultimate pairs present throughout, holding the set fixed leaves it, smaller by whatever the entrants supplied. A cohort is fixed here when its cells carry at least one observed route in both endpoint years. One cohort serves both weightings, and the episode and dollar estimates therefore restrict the same units.

Over 543 matched days on the exact two-leg sample, the pooled change is +25.42 pp by route count and +43.87 pp by value. Restricting to the 39,748 ultimate pairs active in both years leaves +14.85 pp, with a standard error of 1.02 pp, and +39.36 pp, with a standard error of 2.18 pp. Those ultimate pairs carry 55.5% and 57.7% of the two years’ routes and 82.2% and 77.1% of their dollars. Most of the dollar rotation and more than half of the count rotation are therefore observed inside ultimate pairs present on both sides of the comparison. The identity’s within-ultimate-pair term of -0.1 pp by count, whose matched-cell interval reaches only +1.7 pp, then places the surviving rotation: activity moves across established ultimate pairs toward the ones already routing through a stablecoin.

The intermediary menu inside an ultimate pair turns over as well, and a second cohort holds it fixed. Requiring the ultimate pair and the intermediating vehicle to appear in both endpoint years leaves 40,368 ultimate-pair-vehicle cells, a change of +12.94 pp by count and +37.24 pp by value, and 80.4% and 66.1% of the two years’ dollars. Vehicles that intermediate an ultimate pair in one year only thus account for a small part of the measured rotation. The weight moved onto vehicles already available where they were already used.

A third cohort fixes the sequence of exchanges a route ran through and supplies a bound. That sequence is selected jointly with the intermediary. Conditioning on it absorbs part of the rotation’s own content, so the estimate measures a selected comparison. Inside it the count change is -0.18 pp, with a 95% interval of -2.13 pp to $+1.76$ pp, against a pooled +25.42 pp; the value change is +21.74 pp. Episode-weighted rotation is thus bounded to a small fraction of its pooled size once the realised exchange architecture is held constant, while dollar-weighted rotation keeps about half of its own. The cohort retains 37.3% of 2026 dollars against 77.1% for the ultimate pair cohort, and coverage explains part of that gap. [Amiti et al. \(2022\)](#) separate the selection of firms into a currency from the effect of the currency they chose by entering proxies for the choice alongside the

choice itself, and read which of the two dominates. The exchange sequence plays the proxy’s role here, and it dominates the episode measure and shares the dollar measure. Three conditions this comparison would also impose stay outside it: the pools reachable when the route was assembled, the size of the trade, and the best alternative within reach. Each needs pre-transaction state that the route records here do not contain, so none can be estimated in this comparison.

3.3 USDT’s intermediary role

USDC and USDT together account for 92.1% of the increase in stablecoin count share. We examine whether USDT’s changing role is specific to intermediation by comparing its share of intermediary use with its share of appearances at observed route endpoints. Direct routes contribute to the endpoint denominator; longer routes contribute once for each intermediary they use. A ratio above one means that the currency appears more often as an intermediary than at either end of an observed pool route.

On matched January–June dates, the difference between USDT’s intermediary and route-endpoint shares rises from -7.13 pp to $+8.14$ pp. The change is $+15.27$ pp. The ratio of the same two roles, shown as a descriptive full-period comparison in Table 7, rises from 1.06 in 2024 to 1.23 in January–June 2026 by count; the dollar-weighted ratio moves from 0.59 to 1.42. The matched-window gap identifies the transition estimate.

The expansion of cross-exchange routes accounts for 6.3% of USDT’s count-share increase and 14.4% of its dollar-weighted increase. The remaining 93.7% by count and 85.6% by value occurs within the single- and cross-exchange groups. This decomposition conditions on the route that was executed; it does not reveal every alternative available to the router.

The rise in USDT’s intermediary use relative to its observed route-endpoint use could reflect issuance, redemption access, venue coverage, or the distribution of liquidity across pools, and some of these variables may themselves respond to routed activity. Section 5 develops tests that can distinguish these explanations.

Table 7: USDT intermediary use relative to observed route-endpoint use

	2024	2026
Count excess-use ratio (2024 full year; 2026 January–June)	1.06	1.23
Value-weighted excess-use ratio (2024 full year; 2026 January–June)	0.59	1.42
Paired January–June intermediary minus route-endpoint share [pp]	-7.13	$+8.14$

Note: Intermediary and observed route-endpoint shares are formed across native, staked-native, stable, and imported currencies on complete routes that do not return to their starting asset. Direct routes enter the route-endpoint denominator. Excess use is intermediary share divided by route-endpoint share; the gap is their difference. Ratios use the full 2024 calendar year and January–June 2026, while the share gap uses matched January–June dates. Dollar-weighted measures require source, intermediary, and destination values to agree within 20%.

3.4 The cross-section within a day

Every comparison so far reads across years, which invites the reading that the calendar is doing the work. If the vehicle role is simply whatever the sample period happens to do, then the rotation is a description of 2020 through 2026 and not of how an intermediary is chosen. The construct itself does not require the calendar. Excess use compares two roles of the same asset on the same day, so it can be asked entirely within a day, of the assets that were available that day, with the year absorbed into a fixed effect.

We therefore write Equation (3) as the panel regression

$$I_{a,t} = \alpha + \beta D_{a,t} + \mathbf{A}_a' \theta + \mu_t + \eta_a + \varepsilon_{a,t}, \quad (10)$$

where $I_{a,t}$ is currency a 's share of the day's intermediary episodes in pp, $D_{a,t}$ is its share of the day's route-endpoint episodes in the same units, and $\mathbf{A}_a$ contains the asset-class indicators. The specification ladder first estimates class differences without controls, then absorbs date effects μ_t , adds own endpoint demand, and finally replaces the class indicators by currency effects η_a . The slope on endpoint share says whether intermediary use rises more or less than one for one with an asset's own trade demand, and the class intercepts say how much of the role an asset class takes beyond what its demand implies. [Amiti et al. \(2022\)](#) identify currency choice the same way, absorbing the macroeconomic environment in fixed effects so that what remains is the differential response of firms facing the same conditions. The ratio remains the primary extent measure; the regression is the same comparison with the calendar demoted from the identifying variation to a control. The sample is the 1,828,862 currency-days on which 304,572 currencies were quoted at a route endpoint with at least twenty route units, spanning 2,259 days. Inference clusters on date and currency, and on the few-currency cuts on date alone with a thirty-day Bartlett lag, since there the repeated sampling unit is the day.

Table 8 reports the specification ladder in Equation (10); Figure 2 compares the resulting class and currency contrasts in Panels A through C.

Columns (1) and (2) of Table 8 show that absorbing the calendar changes almost nothing. The native premium moves from +34.56 pp pooled to +34.55 pp within day, and the stablecoin premium from +2.42 pp to +2.42 pp. The differences between asset classes are cross-sectional after absorbing period effects. Column (3) conditions on the currency's own trade demand. The native intercept falls to +0.05 pp, while the stablecoin intercept holds at +0.85 pp. Own demand accounts for the native asset's intermediary premium, while the stablecoin premium remains. That is the comparison the search-theoretic account of a medium of exchange asks for, since in [Kiyotaki and Wright \(1989\)](#) an asset is used to intermediate because of what the other assets available at that moment can do, not because of its own turnover. Panel A of Figure 2 displays the class estimates against the residual bucket.

The Own endpoint-demand share row of Table 8 puts the pass-through of demand into the intermediary role at 1.587 across currencies within a day in column (3) and 0.985 within a currency

and within a day in column (4) once 306,830 currency and day effects are absorbed. A currency that gains a point of trade demand gains about a point of intermediary use; the excess above one in the cross-section is composition between currencies.

Restricting the cross-section to the five currencies that carry the vehicle role sharpens it. With the stablecoin class as the base, the native asset intermediates -17.45 pp relative to a stablecoin on the same day and at the same trade demand, with an interval of $[-32.41, -2.49]$ over 11,233 currency-days. Over all 37 classified currencies the same contrast is -1.51 pp, with an interval of $[-4.09, +1.06]$. Panel B of Figure 2 displays both estimates. The two magnitudes answer one question at two scales, and quoting either alone would misstate its size: the vehicle set concentrates the day’s share mass in five currencies, so its coefficients are large, and the classified set spreads the same question across a long tail whose currencies hold almost none of it, so its coefficients are small and the contrast is not separable there. The intermediary role also varies within the stablecoin class. Against a fiat-reserve stablecoin on the same day and at the same demand, an on-chain-collateralised stablecoin intermediates -0.77 pp ($[-1.35, -0.20]$) and a synthetic one -0.61 pp ($[-1.11, -0.11]$), while USDT intermediates -5.32 pp relative to USDC ($[-7.14, -3.51]$). Panel C reports these within-stablecoin contrasts. The class gaining the role remains heterogeneous.

Three features of the design have to be measured before the estimates can be read economically. First, both variables are shares of the same day’s route universe, and a currency cannot intermediate a route it terminates, so a mechanical crowd-out channel exists; it is negligible here, since the mean currency-day uses 0.0008 of the routes it does not itself terminate and only 0.08% of currency-days exceed half of that ceiling, and in any case crowd-out can only push the slope down. Second, each currency’s own episodes enter the denominator of both of its shares. Recomputing both against the day’s episodes excluding the currency’s own raises the slope to 3.209 (0.090), which excludes a shared-denominator explanation. Third, the twenty-unit floor is a choice, and the slope is 1.579 at five units and 1.595 at one hundred. Calendar time returns here only as a sample split: refitting on each half of the sample gives slopes of 1.605 and 1.554, and native intercepts of $+0.70$ pp $[-1.22, +2.63]$ and -0.72 pp $[-2.56, +1.12]$, both intervals covering zero. One result changes under reweighting. Weighting currency-days by route units turns the native intercept to -15.91 pp, so its sign depends on whether each currency-day or each route counts, and we read it as descriptive.

Table 8: Intermediary role within a day

	Intermediary episode share ($I_{a,t}$) [pp]			
	(1) Pooled	(2) Date FE	(3) + demand	(4) + currency FE
Native currency	+34.56 (23.04)	+34.55 (23.04)	+0.05 (0.56)	—
Stablecoin	+2.42* (1.32)	+2.42* (1.32)	+0.85* (0.50)	—
Own endpoint-demand share ($D_{a,t}$)	— —	— —	+1.59*** (0.03)	+0.98*** (0.30)
R^2	0.485	0.485	0.953	0.208
Observations	1,828,862	1,828,862	1,828,862	1,828,862
Currencies	304,572	304,572	304,572	304,572
Dates	2,259	2,259	2,259	2,259
Date fixed effects	No	Yes	Yes	Yes
Currency fixed effects	No	No	No	Yes
Date and currency clustered SE	Yes	Yes	Yes	Yes

Note: Currency-day observations on which the currency is quoted at a route endpoint and the day supplies at least twenty route units: 1,828,862 currency-days over 304,572 currencies and 2,259 days. The outcome $I_{a,t}$ is currency a 's intermediary episode count divided by all intermediary episodes on date t , in pp. The demand regressor $D_{a,t}$ is the currency's appearances as an ultimate-trade endpoint divided by all route-endpoint appearances that day, also in pp; its coefficient is therefore pp of intermediary share per pp of endpoint-demand share. Class premiums are pp differences against the residual unclassified bucket, with standard errors in parentheses clustered on date and currency. The final column absorbs currency and date effects jointly, which leaves class premiums unidentified. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively.

3.5 Venue span and route complexity

Cross-exchange routing expands sharply during the sample. Access to pools on several exchanges can favour an intermediary with liquidity dispersed across venues, while also making direct execution easier to find. Cross-exchange routes grow from 2.0% of route count in 2020 to 57.2% in 2026 and account for 79.1% of dollar-weighted route value in 2026.

Intermediation becomes less common over 2022–2026 even as cross-exchange routing grows. Its count share falls from 13.95% to 11.92% across the full exchange set and to 8.60% among five exchanges observed throughout the comparison. The changes are -2.03 pp and -5.35 pp. Stablecoin share nevertheless rises within every observed group defined by exchange span and route complexity. Its count-share increase is $+7.45$ pp larger across exchanges than within one exchange; the dollar-weighted estimates rise in both groups. The Intermediary episodes and Routed value panels of Figure 3 report these changes by single- and cross-exchange route scope.


Those estimates compare two endpoint years. The calendar supplies the variation that identifies them. A second reading of the same stratification removes that reliance. Each trading day contributes one observation for its single-exchange routes and one for its cross-exchange routes, and a day effect absorbs everything the two halves share, from gas prices to market direction. The remaining difference measures how far apart the two halves of one day’s route universe stand.

Cross-exchange routes carry a larger stablecoin share of routed value. Over 2,234 days on which both halves are measured, their value-weighted share exceeds that of single-exchange routes by $+14.18$ pp, with a 95% interval of $[+7.0, +21.4]$; weighting each day-group cell by its dollars returns $+12.12$ pp. The route-count estimate is $+3.73$ pp, with an interval of $[-2.2, +9.7]$. Exchange span therefore tracks which currency carries the value that moves through an intermediary, while the same-day evidence remains descriptive about how often an indirect route reaches for a stablecoin at all. [Chen and Duffie \(2021\)](#) model trade split across exchanges as thinning depth at each one while improving allocation overall. Thin depth at a single venue binds first on the largest orders, and a difference that opens in routed dollars while route counts stay level fits that reading.

The dollar gap is also unsteady across the sample. Estimated year by year, it opens at $+69.7$ pp in 2020, closes to -3.2 pp in 2023 on an interval of $[-9.2, +2.8]$, and reopens to $+14.5$ pp in 2026. A design that reads only the endpoint years of the sample would record a gap at both ends and miss the closure between them. Route exchange span is selected, so these differences describe the composition of two selected populations. The effect of widening a route’s venue set remains outside this comparison.


Taken together, the estimates separate cross-exchange trading, the incidence of intermediation, and the composition of intermediaries. Wider venue span coincides with less intermediation overall, while stablecoins gain among the indirect routes that remain in both single- and cross-exchange groups. Section 5.1 turns that stratification into a test of whether the aggregate rotation is a by-product of how routes were built. Isolating the effect of router access requires comparisons among the same ordered source–destination asset ultimate pairs, trade sizes, and feasible paths.

Figure 2: Within-day intermediary role, conditional on own trade demand



Note: Each row is one coefficient from a regression of a currency’s share of the day’s intermediary episodes on its own share of that day’s route-endpoint episodes, both in pp, with date fixed effects absorbing the calendar. Bars are 95% intervals. A filled marker means the interval clears zero and an open marker means it includes zero; rows are drawn at the same weight either way. Panel A measures each class against the residual unclassified bucket over every endpoint-supported currency-day. Panel B measures the native asset against a stablecoin on the five named vehicles and on all classified currencies; the two samples put different amounts of the day’s share mass behind the same indicator, which makes the magnitudes incomparable. Panel C measures backing regimes against a fiat-reserve stablecoin and USDT against USDC inside the stablecoin class. The three panels use separate horizontal scales.

Figure 3: Stablecoin intermediary share rises within single- and cross-exchange routes



Note: The sample contains complete routes that do not return to their starting asset, including routes with more than two legs. Each route contributes once for every intermediary it uses. Stablecoin shares use all intermediary types as the denominator. Count and dollar-weighted estimates are reported separately on matched January–June dates in 2024 and 2026. Dollar-weighted measures require source, intermediary, and destination values to agree within 20%.

4 Comparing direct and intermediated execution

Realised routes identify the assets that carried each trade. Execution cost is a separate object: it compares those routes with alternatives available immediately beforehand, quoted at the same trade size and over a common set of exchanges. We hold the source i , destination o , dollar input q , and pre-transaction pool state e^- fixed. The best direct and two-leg outputs through intermediary k are

$$O_{i,o,q,e}^D = P_{o,e} \max_{p \in \mathcal{P}_{i,o,e^-}} Q_p(i \rightarrow o, \Delta_i(q, e); \chi_{p,e^-}),$$

$$O_{i,o,k,q,e}^I = P_{o,e} \max_{p' \in \mathcal{P}_{k,o,e^-}} Q_{p'}\left(k \rightarrow o, \max_{p \in \mathcal{P}_{i,k,e^-}} Q_p(i \rightarrow k, \Delta_i(q, e); \chi_{p,e^-}); \chi_{p',e^-}\right)$$

This comparison extends the realised-versus-optimised allocation benchmark of [Xi and Moallemi \(2026\)](#) from parallel pools for one atomic asset pair to sequential routes through an intermediary.

The vehicle-route shortfall relative to direct execution is

$$\Delta C_{i,o,k,q,e}^D = \frac{O_{i,o,q,e}^D - O_{i,o,k,q,e}^I}{O_{i,o,q,e}^D},$$

reported in basis points as $10^4 \Delta C_{i,o,k,q,e}^D$. Positive values mean that the direct route returns more output before gas. Swaps and liquidity changes are replayed in on-chain order because a daily or hourly closing state incorporates events that occurred after the trade. The observed event marks settlement; a quote formed earlier can reflect a different opportunity set.

We compare hypothetical legs only within a common executable range. Let $\tilde{Q}_\ell$ denote leg ℓ 's output at the pool's marginal rate. Each leg satisfies

$$\pi_\ell(\Delta_i, e^-) = \frac{\tilde{Q}_\ell(\Delta_i, e^-) - Q_\ell(\Delta_i, e^-)}{\tilde{Q}_\ell(\Delta_i, e^-)} \leq \bar{\pi}, \quad \bar{\pi} = 0.05. \quad (11)$$

The 5% own-price-impact ceiling applies across pool families; contract-specific limits add further restrictions. Appendix B gives the pricing equations and Appendix D calibrates the support rule.

Receipt gas measures the fixed execution toll attached to adding a vehicle leg. That toll matters because gas fees can change both trading costs and liquidity concentration in decentralised exchange ([Caparros et al., 2024](#)). In the executed-route gas sample, a stable-vehicle route in 2024 uses 94,290 more gas than the median direct route, 50.2% of the direct-route median. Evaluated at that year's median gas price and WETH price, the extra toll is \$3.24, so a trade needs about \$32.4k of notional before the extra gas is below one basis point. In 2026 the median stable-vehicle route uses 233,640 more gas than the direct-route median, but the dollar toll falls to \$0.08 and the one-basis-point notional falls to \$762. Pricing the same fixed toll at each stable-vehicle route's own gas price sharpens the interpretation: median stable-vehicle notional falls from \$832 to \$82, while the median fixed toll falls from 44.3 bp to 16.2 bp and the share of stable-vehicle routes with a fixed toll below 25 basis points rises from 39.5% to 55.7%, the trade-size version of the fixed-cost channel

in [Caparros et al. \(2024\)](#). Thus the extra-hop overhead remains visible in gas units while becoming economically small for ordinary routed notionals. This evidence supports the interpretation that vehicle formation spread to routes for which fixed execution overhead had ceased to be a large barrier. The estimand is the executed route’s fixed gas toll; a stable-versus-WETH all-in price ranking remains outside this check.

The benchmark represents public on-chain liquidity in the declared exchange set. Private order flow, financing costs, losses from reverted transactions, and pools outside that set can alter a router’s choice. Pool-family coverage is economically material. In the sampled Curve exercise, pools outside the StableSwap model carry 34.2% of assessed volume, including 65.2% of native-leg volume and 21.1% of stablecoin-leg volume. Hook-bearing or dynamic-fee Uniswap v4 pools, SushiSwap v3, and several Balancer families also use pricing states outside the benchmark. [Appendix E](#) quantifies these boundaries.

5 What can account for the rotation?

The allocation results sharpen the mechanism question. Stablecoin dominance rose chiefly as routed activity moved across source–destination ultimate pairs and the incidence of intermediation changed, while net replacement of a native intermediary within continuing ultimate pairs contributed little. Any explanation for the rotation must therefore account for the assets being exchanged as well as the intermediary carrying the route.

5.1 Path complexity and exchange span

A competing account assigns the rotation to how routes are assembled. Paths grew longer and reached across more exchanges over the sample, and long paths already favoured stablecoins: in 2024 stablecoins carried 44.5% of native-plus-stable intermediary episodes on routes of more than two legs, against 16.9% on two-leg routes. That account predicts a roughly flat stablecoin share inside each complexity class and an aggregate increase produced by activity moving toward longer routes that were already stablecoin-intensive.

Stablecoin share rises within every complexity class. Among intermediary episodes on two-leg routes, the stablecoin share rises +25.4 pp by count, with a standard error of 1.05 pp, and +43.9 pp by value (2.02 pp). Among episodes on routes of more than two legs it rises +16.3 pp (1.86 pp) and +35.9 pp (1.68 pp), from a base of 44.5% to 60.7% by count. Both classes move, and the class with the shortest paths moves further, which leaves reweighting toward complex routes too small to supply the aggregate change. Dividing each complexity class again by whether the route touched one exchange or several leaves four groups, and every one of them rises; the weakest is +17.0 pp by count (1.36 pp) and +25.6 pp by value (1.52 pp). The stablecoin gain is thus common to short and long paths and to routes that stay on one exchange and cross-exchange routes.

Two boundaries limit what this comparison licenses. The number of legs records the path a trade actually took, so it measures construction; extra legs have no ordinal interpretation for execution

quality. Complexity and exchange span are also jointly selected outcomes, since the same conditions that make a longer or wider path attractive can bear on the choice of intermediary. The comparison locates the stablecoin increase within each route-complexity and exchange-span class. The effects of routing software and exchange span remain unidentified.

5.2 The software that assembles the path

Traders on these exchanges seldom assemble a path by hand. A router contract holds the pool set, searches it, and returns a sequence of legs, and the routers in widest use over this sample were rewritten three times: the Auto Router in September 2021, its cross-version successor that December, and the Universal Router in November 2022.⁵ Each release widened the set of paths an ordinary trader could reach without extra effort. The alternative this raises is that the intermediary is a residue of the search: the software found a path through a stablecoin because such paths existed and had become cheap to look for, so the rotation would record an improvement in search and not a change in what traders were willing to hold between legs. Search does leave room to move. On parallel pools for a single atomic pair, realised routes fall measurably short of the optimum available in the same pool state (Xi and Moallemi, 2026), so a router that closes part of that gap changes which paths trades actually take.

The market-microstructure design used for staggered trading technologies is unavailable here. That literature identifies the effect of an exchange technology when the exchange grants it to some securities before others, because the securities still waiting absorb whatever else is moving in the market at the same time. A router release admits no such comparison. The contract is deployed once and is available immediately for every atomic pair, at every exchange the router indexes, to every trader, so there is no untreated group and no difference to take. What a symmetric window around each release can measure is whether the structure of routes moved at all when the search over them was rewritten, and that weaker object is the one the alternative needs. If automated search produced the intermediary, the incidence of intermediation should step up when the search was automated.

Route composition remains stable around the releases. Table 9 reports route composition over the 60 days on either side of each release, excluding the release date itself, on the daily calendar used throughout the paper. The share of economic routes carrying a third asset between the endpoints falls by -5.7 pp at the first release, rises by $+0.8$ pp at the second, and falls by -0.7 pp at the third, moving from 20.8% before the first release to 12.5% after the last. The largest increase at any of the three is $+0.8$ pp. Path length is as still: the mean number of legs on an indirect route moves by at most 0.049 at any release, and the share of indirect routes running past two legs changes by a few pp in either direction. Indirect routing as a whole behaves the same way, falling sharply at the first release and moving by less than half a pp at the other two.

⁵Uniswap Labs, *Introducing Auto Router*, September 16, 2021, and *Introducing Permit2 and Universal Router*, November 17, 2022. The cross-version release is dated from Uniswap Labs' *v2-versus-v3 liquidity-provider study* of December 16, 2021.

Table 9: Route structure around public router releases

	Indirect	Intermediated	Cross-exchange	Legs	Exchanges	Over two legs
<i>Panel A: Auto Router, September 16, 2021 (60 days each side)</i>						
Before	22.7	20.8	8.5	2.08	1.10	6.6
After	18.9	15.1	11.0	2.10	1.11	8.1
Change	-3.7	-5.7	+2.5	+0.01	+0.01	+1.5
<i>Panel B: Cross-version Auto Router, December 16, 2021 (60 days each side)</i>						
Before	18.5	13.9	11.3	2.11	1.11	8.4
After	18.6	14.8	14.4	2.16	1.16	11.3
Change	+0.2	+0.8	+3.1	+0.05	+0.05	+2.9
<i>Panel C: Universal Router, November 17, 2022 (60 days each side)</i>						
Before	19.8	13.2	26.4	2.19	1.22	13.1
After	19.5	12.5	24.4	2.20	1.20	12.0
Change	-0.3	-0.7	-2.0	+0.01	-0.03	-1.1

Note: Symmetric windows of 60 calendar days either side of each release, excluding the release date. Indirect is the share of economic routes with more than one leg; Intermediated is the share carrying a third asset between the endpoints; Cross-exchange is the share of intermediated routes touching more than one exchange; Over two legs is the share of indirect routes with more than two legs; Legs and Exchanges are means per indirect route. Shares and their changes are in pp. Economic routes exclude canonical endpoint round trips. The windows are market-wide and carry no split by intermediary currency. A release is available to every trader at once, so these are descriptive route compositions.

One margin does move, and it is the margin a router should move. Among intermediated routes, the share touching more than one exchange rises by +2.5 pp at the first release from a base of 8.5% and by +3.1 pp at the second; at the third it changes by -2.0 pp, from a level near a quarter of intermediated routes to 24.4%. Automated search integrated the exchange set without sending trades through more assets. The distinction is the one that governs this paper’s measure, which counts the asset standing between the endpoints and not the number of venues its legs visit, and it is the same separation the exchange-span stratification of Section 5.1 makes within a single date.

Router releases expand exchange span while intermediary incidence remains stable. Three features of the windows bound that comparison. The first two windows overlap: the later window of the first release and the earlier window of the second share thirty days, so the first two comparisons are drawn partly from the same trading. The third window also contains a major market event. The earlier window of the third release contains the November 2022 failure of a large centralised exchange, and the composition change recorded there belongs at least as much to that event as to the router. A release date, finally, marks when a contract became callable while trader adoption can begin later, so a window straddling gradual adoption understates a movement that arrived slowly.

5.3 The pricing rule at each leg

The third alternative centers on the exchanges that price each leg, because AMMs apply different pricing rules across atomic pairs. The constant-product rule sets a price from pool reserves alone

and moves it with trade size relative to depth, so it charges a trade for the price impact it causes regardless of how closely the two assets substitute. Curve’s invariant instead interpolates toward a flat curve in the neighbourhood of parity, which is exactly where two dollar-pegged assets trade, and Appendix B.3 states the interpolation the venue applies. On this account the rotation is an adoption of technology: a leg through a stablecoin became cheap once venues priced pegged assets on a curve built for them, and the vehicle role passed to stablecoins because of where their legs could be priced, so no monetary property of the asset need enter the account at all. The timing is favourable to it, since stable-specialised pools grew over the same years as the rotation.

The account can be tested by restricting attention to routes on which the pricing rule never changed. A route component enters the principal panel’s constant-product scope only when every one of its legs executes at Uniswap v2 through v4 or at SushiSwap v2 or v3, each of which prices locally on the product of reserves, concentrated liquidity inside an active tick included. No leg inside that scope can have been priced by a curve specialised for pegged assets, so a technology account predicts that stablecoins gain nothing there. The quantity compared is the excess-use ratio of Section 3.3, an asset’s share of intermediary activity divided by its share of route-endpoint demand, formed separately inside each scope, so that the denominator for a venue family is the endpoint demand observed in that same family.

The constant-product restriction produces an even larger rotation. Where every leg is priced by the reserve product, stablecoin excess use rises from 0.78 in 2024 to 1.33 in 2026 by value, a movement of +0.55 against +0.40 over all venues, and the ratio crosses one inside the restricted scope in the same comparison. Native assets fall to 0.62 there over the same years. By count, where stablecoins were already over-represented as intermediaries at the start of the sample, the restricted and unrestricted paths are almost the same: 1.28 to 1.42 inside the scope against 1.27 to 1.41 over all venues. The scope also carries substantial activity, since routes priced end to end by the constant-product rule carry 84.8% of vehicle-currency intermediary episodes in 2026. Table 10 reports every year and every venue family.

The venue that embodies the alternative supplies no intermediation at all. In each year of the sample, route components all of whose legs execute on Curve contain no intermediary episode whatever. The stable-specialised curve is used to exchange two pegged assets with one another, not as a place through which a route passes on its way somewhere else, so the pricing rule that makes stablecoin legs cheap operates on the direct trades that the vehicle measure excludes by construction. Curve legs do appear inside routes that also touch another exchange, and those routes are counted in the all-venue scope; what the sample contains no instance of is a route intermediated entirely within the specialised invariant. Balancer, whose pools mix weighted and stable families and which the panel therefore leaves technologically unlabelled, carries 39,536 intermediary episodes in its busiest year against 8,767,213 in the busiest year of the constant-product scope, and its ratios swing accordingly; it is reported solely to delimit the sample’s composition.

What the comparison licenses is again bounded. Holding the pricing rule fixed supplies only a descriptive restriction: venues are chosen alongside routes, and a trader who would have used

Table 10: Vehicle excess use by venue pricing family

	All venues		Constant product		Balancer		Curve	
Year	Count	Value	Count	Value	Count	Value	Count	Value
<i>Panel A: Stablecoins</i>								
2020	1.45	0.42	1.50	0.42	no routes		no intermediation	
2021	1.18	0.89	1.20	0.92	2.09	1.43	no intermediation	
2022	1.46	1.08	1.50	1.12	1.91	1.00	no intermediation	
2023	1.89	0.94	1.95	1.00	1.46	0.37	no intermediation	
2024	1.27	0.84	1.28	0.78	1.01	0.27	no intermediation	
2025	1.23	1.21	1.23	1.20	0.70	0.09	no intermediation	
2026	1.41	1.24	1.42	1.33	1.75	1.68	no intermediation	
<i>Panel B: Native asset</i>								
2020	0.91	1.56	0.90	1.37	no routes		no intermediation	
2021	0.94	1.13	0.94	1.09	0.24	0.66	no intermediation	
2022	0.85	0.96	0.85	0.93	0.35	0.90	no intermediation	
2023	0.87	1.16	0.88	1.08	0.50	0.53	no intermediation	
2024	0.95	1.20	0.95	1.29	0.22	0.20	no intermediation	
2025	0.89	0.67	0.92	0.85	0.19	0.19	no intermediation	
2026	0.77	0.58	0.80	0.62	0.51	1.30	no intermediation	

Note: Excess use is an asset type’s share of intermediary activity divided by its share of route-endpoint demand, formed over the vehicle currencies within each scope. A route component enters a scope only when every one of its legs executes at a venue in that family. The principal panel’s constant-product family is Uniswap v2 through v4 and SushiSwap v2 and v3. Uniswap v1 enters the separate forced-routing architecture sample. Balancer is listed separately and left technologically unlabelled because its pools mix weighted and stable families. Value ratios use all routes; the 20% value-agreement support of Table 2 is narrower. Balancer has no route components before 2021, and 2026 covers January through June.

a stable-specialised pool may appear in the constant-product scope only because that pool was unavailable for the atomic pair at hand. The ratio’s denominator removes the most obvious composition problem, since each scope is judged against the endpoint demand observed inside it, but selection into scopes remains. The rotation therefore persists beyond stable-specialised pricing: it is present, large, and in the same direction where that pricing was never applied.

5.4 Liquidity in the vehicle’s own pools

The classical vehicle-currency mechanism links trading volume to lower transaction costs (Krugman, 1980). On a decentralised exchange, the relevant markets are the pools connecting an intermediary to other assets. Trading can attract provider capital, deeper pools can lower price impact, and lower price impact can draw further trading. Pool size remains an equilibrium outcome because providers balance fee income against adverse-selection risk and the cost of adjusting their positions (Lehar and Parlour, 2025). Deposited capital measures funds committed to a pool; a size-specific quote measures how much liquidity is executable within a stated price impact. The distinction matters. Concentrated-liquidity providers can alter executable depth by changing where capital is placed along the price curve, so deposited capital and executable liquidity coincide only under particular pool designs (Adams et al., 2021).

That mechanism operates through the two atomic-pair venues joining an intermediary to the ultimate endpoints. Deeper stablecoin pools can lower the price impact of carrying an ultimate trade through a stablecoin, so the classical account asks whether flow moves away from the incumbent intermediary after a stable bridge becomes feasible. Section 3.2 conditions on realised routes in both endpoint years: within those matched ultimate pairs, stablecoin share changes by +0.2 pp by count, with an interval that excludes increases beyond +1.7 pp. Bridge feasibility and adoption remain pooled in that estimate. A separate event design therefore dates the feasible-set expansion before asking how the realised route changes.

Let event e identify an ordered ultimate pair (i, j) , route scope, and the first day T_e on which DAI, USDC, or USDT has positive prior-calendar deposited capital on both atomic pairs, with two-leg support present on at least 24 of the next 30 calendar days. The sample contains ultimate pairs that traded through WETH before T_e , have no earlier observed stablecoin route, and trade on at least three days in the prior month and each follow-up window. For route counts and supported value $q \in \{R, V\}$, the paired change is

$$\Delta S_{e,h}^q = S_{e,h}^q - S_{e,[-30,-1]}^q,$$

where h is either the first 30 days or days 30 through 119 after bridge establishment. Share changes use the harmonic mean of pre- and post-event mass; route-intensity changes weight each event equally. Standard errors cluster by ordered ultimate pair and event date.

Panel A of Table 11 records the change around first persistent support. Across 865 events, stable route share rises by +7.93 pp in the first month and supported-value share by +3.11 pp. Native-route intensity changes by -0.20 log points, while total route intensity changes by -0.10 log points. The count-share increase over days 30 through 119 is +8.68 pp; the supported-value estimate is +6.94 pp. These event-window estimates date stable support and the associated route allocation. The observed support event is an equilibrium outcome.

The continuous-depth design measures the capital available on both routes after that support appears. Let $K_{av,t-1}$ be exact prior-calendar deposited capital in the atomic pair joining endpoint a to vehicle v , and let $\mathcal{S} = \{\text{DAI, USDC, USDT}\}$. Define the best stable weak-leg depth D_{et}^S , WETH weak-leg depth D_{et}^W , their stable relative-depth share Q_{et} , and stable route-count share S_{et}^R by

$$\begin{aligned} D_{et}^S &= \max_{v \in \mathcal{S}} \min\{K_{iv,t-1}, K_{vj,t-1}\}, & D_{et}^W &= \min\{K_{iW,t-1}, K_{Wj,t-1}\}, \\ Q_{et} &= \frac{D_{et}^S}{D_{et}^S + D_{et}^W}, & S_{et}^R &= \beta Q_{et} + \alpha_e + \lambda_{m(t)} + \eta_{a(t)} + \varepsilon_{et}. \end{aligned} \tag{12}$$

Here W denotes WETH, $m(t)$ is calendar month, and $a(t)$ groups event age into seven-day bins. Equation (12) absorbs bridge-event fixed effects, weights by the event day's WETH-plus-stable route count, and clusters standard errors by ordered ultimate pair and date.

Panel B of Table 11 reports the continuous-depth result. The first-month sample contains 11,327 comparable ultimate-pair days; on 7,752 of them, the best stable weak-leg depth is below one-tenth

of WETH weak-leg depth, and stablecoins carry 2.4% of route counts in that range. Within a bridge event, a 1 pp increase in Q_{et} is associated with +0.64 pp higher stable route share. Stablecoins carry 53.0% of routes when their weak-leg depth is at least WETH depth, and 69.9% when it is at least twice WETH depth. Deposited depth is therefore the proximate liquidity-formation margin connecting support to route allocation. The design identifies a relationship with local capital stocks; executable quotes and causal provider-supply shifts require separate measurement and design.

Table 11: Stable-bridge establishment, relative depth, and route allocation

Outcome	First 30 days	Days 30–119
<i>Panel A. Changes around first persistent stable support</i>		
	+7.93***	+8.68***
Stable route share [pp]	(1.88)	(2.57)
	+3.11**	+6.94
Stable supported-value share [pp]	(1.53)	(4.53)
	−0.20***	−0.39***
$\log(1 + \text{native routes per active day})$	(0.04)	(0.06)
	−0.10**	−0.29***
$\log(1 + \text{total routes per active day})$	(0.04)	(0.05)
	−0.17	−0.61***
$\log(1 + \text{native supported value per active day})$	(0.11)	(0.19)
	+0.06	−0.32*
$\log(1 + \text{total supported value per active day})$	(0.10)	(0.18)
Bridge events	865	818
<i>Panel B. Stable-bridge competitiveness relative to WETH</i>		
	+0.64***	+0.75***
Relative-depth slope [pp per +1 pp]	(0.07)	(0.05)
Stable route share, depth < 0.1× WETH [%]	2.4	2.4
	53.0	61.2
Stable route share, depth ≥ WETH [%]	(4.6)	(13.2)
	69.9	70.7
Stable route share, depth ≥ 2× WETH [%]	(2.3)	(9.3)
Active ordered-ultimate-pair days	11,327	17,027

Note: The unit is an ordered ultimate pair and route scope around the first persistent DAI, USDC, or USDT bridge. Support requires positive prior-calendar deposited capital on both atomic pairs in Uniswap V2 or SushiSwap V2 on at least 24 of the next 30 calendar days. Every ultimate pair has prior WETH-mediated routes and no earlier observed stablecoin route. In Panel A, the pre-period is the 30 calendar days before support; share rows report paired pp changes using harmonic pre/post route or supported-value mass, and intensity rows report paired log changes with equal event weights. Panel B compares the best stablecoin weak-leg capital with WETH weak-leg capital on post-event active days. The relative-depth slope is Equation (12); threshold rows are route-count-weighted means. Parentheses contain standard errors two-way clustered by ordered ultimate pair and date. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively. Deposited capital measures the prior-calendar pool stock and supplies a proxy for liquidity formation; executable quotes and causal provider-supply shifts require separate measurement and design.

The local bridge-depth analysis supports that earlier-margin interpretation. Let $c = (i, j, t, s)$ denote an ordered ultimate-pair, date, and route-scope opportunity; let v index vehicles; and define $B_{vc} = \log(1 + \min\{K_{iv,t-1}, K_{vj,t-1}\})$, the prior-calendar deposited capital on the weaker of the two

atomic pairs. The main route-choice specification is

$$R_{vc} = \beta_B B_{vc} + \theta_0 \log(1 + N_{v,t}^{(-c)}) + \theta_1 \log(1 + N_{v,t-30}) + \theta_2 \log(1 + P_{v,t-30}) + \alpha_c + \delta_v + \varepsilon_{vc}, \quad (13)$$

where R_{vc} is the vehicle's share of five-vehicle route mass, $N_{v,t}^{(-c)}$ is same-day leave-one-opportunity-out route reach, and $N_{v,t-30}$ and $P_{v,t-30}$ are prior-30-day route and ultimate-pair reach. The regression weights observations by five-vehicle route count, absorbs opportunity and vehicle effects, and clusters by ordered ultimate pair and date. For each opportunity in which at least two of the five headline vehicles have both atomic legs supported, we rank vehicles by B_{vc} . The deepest supported bridge carries 84.1% of five-vehicle route mass, rising from 75.5% in 2024 to 86.5% in 2026, across 58,483 opportunity cells and 2,923 ordered ultimate pairs. With opportunity and asset effects, a one-log-point larger weaker-leg capital stock is associated with +6.67 pp higher route share, with standard error 1.21 pp; the corresponding stablecoin total slope is +7.73 pp, with standard error 1.46 pp. Adding the reach controls in equation (13) leaves the local-depth coefficient at +6.79 pp, with standard error 1.15 pp; same-day global route reach contributes +9.10 pp, with standard error 2.81 pp. Restricting the same vehicle set to first-observed ultimate pair dates, and retaining zero-depth alternatives as part of the birth choice set, gives the entry margin directly: across 835 vehicle rows in 167 entry opportunities, a one-log-point larger local bridge predicts +5.35 pp higher route share and +5.28 pp higher selection probability, with standard errors 0.51 pp and 0.57 pp. Prior-30-day route reach is estimated at +8.78 pp with standard error 5.72 pp, so local depth adds information beyond existing vehicle popularity. A two-leg bottleneck specification gives the physical-market version of the mechanism: holding the stronger leg and reach fixed leaves the weaker-leg slope at +6.61 pp, with standard error 1.17 pp, while a one-log-point leg imbalance predicts -5.43 pp route share, with standard error 1.14 pp. The local-depth slope also survives all 5 leave-one-vehicle reruns of the same horse race, with the weakest retained slope +5.98 pp and standard error 1.54 pp, which makes the result robust to every single-vehicle omission. Inside the stable-only supported race, 12,048 opportunity cells across 296 ordered ultimate pairs compare DAI, USDC, and USDT directly. With the same opportunity effects, local depth, and reach controls, the 2026 USDC interaction is +34.65 pp with standard error 17.18 pp, the USDT interaction is +55.29 pp with standard error 13.17 pp, and local depth remains +6.01 pp with standard error 1.97 pp. Conditional on all three stablecoin bridges being feasible in the same opportunity, the 2026 allocation is concentrated in USDC and USDT relative to DAI. Deposited capital remains a predictive state variable. Route choice is therefore organised around both vehicle identity and ultimate-pair-specific two-leg depth where ultimate-pair entry and activity growth create scope for vehicle formation.

The USDC and USDT interactions answer a different conditional question: among stablecoin bridges already supported in the same opportunity, which issuer carries more route flow in 2026? They rank feasible stablecoin alternatives after availability and do not date bridge establishment.

Table 12: Dynamic feedback between route use and local bridge depth

Feedback equation	30 days	120 days
$R_{b,t} \rightarrow \Delta B_{b,t+h}$, pooled [log points]	+0.020*** (0.005)	+0.097*** (0.013)
$R_{b,t} \rightarrow \Delta B_{b,t+h}$, stable [log points]	+0.024*** (0.005)	+0.115*** (0.012)
$B_{b,t} \rightarrow \Delta R_{b,t+h}$, pooled [pp]	+1.262*** (0.189)	+1.595*** (0.271)
$B_{b,t} \rightarrow \Delta R_{b,t+h}$, stable [pp]	+1.773*** (0.228)	+2.463*** (0.292)
Rows / ultimate pairs / dates	51,240 / 503 / 303	18,304 / 340 / 123
Asset and date effects		Yes
Two-way clustered covariance		Ordered ultimate pair and date

Note: Each cell reports a predictive coefficient with two-way clustered standard errors in parentheses. For source–vehicle–destination bridge b , $R_{b,t}$ is the vehicle share of five-vehicle route mass and $B_{b,t} = \log(1 + \min\{K_{iv,t}, K_{vj,t}\})$ is deposited capital on the weaker atomic pair; Δ is the exact-horizon value minus the origin-date value. Route-use rows scale $R_{b,t}$ by 10 pp and report $\Delta B_{b,t+h}$ in log points. Bridge-depth rows scale $B_{b,t}$ by one log point and report $\Delta R_{b,t+h}$ in pp. All regressions absorb vehicle and origin-date effects, control for the other current state variable, weight by five-vehicle route count, and cluster by ordered ultimate pair and date. Stable rows add the stablecoin interaction and report the stable total. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively. These local feedback regressions use observational variation.

A dynamic bridge analysis gives the feedback version of the volume–cost mechanism in [Krugman \(1980\)](#). Let b index a source–vehicle–destination bridge, $R_{b,t}$ be the vehicle’s share of five-vehicle route mass in that local opportunity, and $B_{b,t}$ be the log weak-leg bridge depth. The two predictive equations are

$$\begin{aligned}\Delta B_{b,t+h} &= \beta_R R_{b,t} + \gamma_R B_{b,t} + \alpha_c + \delta_t + \varepsilon_{b,t+h}, \\ \Delta R_{b,t+h} &= \beta_B B_{b,t} + \gamma_B R_{b,t} + \alpha_c + \delta_t + \eta_{b,t+h}.\end{aligned}\tag{14}$$

Matching the same bridge at exact future horizons, a 10 pp higher current route share predicts +0.020 log points more bridge-depth growth after 30 days and +0.097 after 120 days. In the other direction, a one-log-point larger current bridge predicts +1.26 pp more route-share growth after 30 days and +1.60 pp after 120 days. The feedback margin is stronger for stable vehicles: the corresponding 120-day stable route-share to depth-growth coefficient is +0.115 for a 10 pp route-share difference, and the stable depth to route-share coefficient is +2.46 pp. These rows absorb vehicle and origin-date effects and cluster by ordered ultimate pair and date. Consistent with constrained and active liquidity provision ([Lehar and Parlour, 2025](#); [Huang et al., 2025](#)), the estimates identify predictive comovement: route use and local bridge depth forecast one another inside the same bridge. One-for-one provider funding lies outside the design.

Table 12 reports the coefficients from Equation (14).

The allocation of capital itself cuts against a flow-chasing interpretation. In the daily five-vehicle V2 panel, WETH is the largest deposited-capital vehicle on 100.0% of 2,239 days, while stablecoins lead excess use on 81.3% of days. From 2024 to 2026, stablecoins’ share of vehicle deposited capital falls from 17.6% to 11.1%, while their share of vehicle intermediary route counts rises from 18.1% to 46.6%, leaving their 2026 route share $4.2\times$ their capital share. On the matched January-through-June daily calendar, the stablecoin route-minus-capital gap expands by +33.5 pp,

while the WETH gap moves by -36.6 pp. Holding the date fixed, and controlling for endpoint share, venue count, and pool count, stablecoin vehicles still carry $+13.3$ pp more route share than deposited-capital share.

The signed gap is associated with capital-share convergence: within vehicle and origin-date effects, a 10 pp route-minus-capital gap predicts $+0.61$ pp higher vehicle capital share after 30 days, and $+1.65$ pp after 120 days. This specification has a mechanical mean-reversion component because origin capital share enters negatively in the gap and again through the origin value subtracted from future capital share. The capital-share and capital-rank rows therefore describe convergence; provider reallocation requires exogenous variation. For stable vehicles, the corresponding 30-day signed-gap coefficient is $+0.78$ pp. The stable association varies by issuer: at 120 days, the same 10 pp gap predicts $+5.17$ pp more DAI capital share, $+0.75$ pp more USDC capital share, and $+0.19$ pp for USDT. In log deposited capital the USDC and USDT associations are $+14.44\%$ and $+6.94\%$. Rank transitions show the same convergence pattern: for stable vehicles, a 10 pp route-minus-capital gap predicts $+0.23$ positions of capital-rank catch-up after 120 days, while route rank moves -0.27 . The pattern is asymmetric. In a piecewise version of the same design, a 10 pp stable capital overhang, meaning capital share exceeds route share by 10 points, predicts a -3.91 pp capital-share change after 30 days and -7.11 pp after 120 days. Aggregating the three stablecoins into a dollar-vehicle basket shows where much of the convergence appears. With activity controls and a 30-day calendar HAC, a 10 pp stable-basket route-capital shortfall predicts $+2.30$ pp higher stable-basket capital share after 120 days and -2.02 pp lower WETH capital share; the WBTC response is -0.28 pp. The extensive-margin association appears in venue reach: for stable vehicles, a 10 pp route-minus-capital gap predicts $+0.56\%$ future log venue-count change after 30 days and $+0.55\%$ after 120 days. The 120-day future log pool-count estimate is -2.73% . Stable capital is also more concentrated: by 2026 its capital-weighted top-pool share is 52.2% , versus 11.4% for non-stable vehicles, and the same 10 pp stable gap predicts $+2.68$ pp higher future pool-capital HHI after 120 days. The V2 stock, Portfolio rank, and Pool footprint rows of Table 13 display the corresponding stock, rank, and incumbent-pool coefficients.

The USDC/SVB stress window gives the same separation in event time. Stable route demand rose sharply, with stable basket route share increasing by $+11.4$ pp inside the five-vehicle set, while stable deposited-capital share moved -2.0 pp during the same window and -3.3 pp afterward. USDC’s own capital share was -3.2 pp lower after the window, while WETH’s was $+3.1$ pp higher. The episode combines peg stress, arbitrage, and demand, so it remains descriptive. It still falsifies a mechanical capital-chases-current-routes story: the route shock is visible immediately, while the V2 capital stock shifts away from stablecoins over the same event horizon.

A same-pool specification gives the rival its best shot. Across pool-vehicle rows, a 10 pp route-minus-capital gap predicts $+2.36\%$ future log deposited-capital growth after 120 days; for stable vehicles the corresponding total is $+0.72\%$. Splitting future exact-state V2 capital by whether the pool was already active at the origin date concentrates the response in incumbent pools: stable shortfalls predict $+12.39\%$ incumbent-pool capital after 120 days; the entrant-pool and total-capital

estimates are +1.06% and +4.39%. Rent incidence gives the same answer in the V3 fee archive: across 208,360 same-pool horizon rows in 3,034 pools, a 10 pp stable route-minus-capital gap predicts +0.10% future log fees and +1.53% future log volume after 120 days. Fee and volume grow in levels; the corresponding estimates for future log fee yield and future log volume turnover are -0.27% and 0.00%. Provider actions respond on a separate margin. In 11,195 vehicle-days built from Uniswap v3 mint and burn logs, a 10 pp stable route-minus-capital gap predicts +67.89% more future mint events and +64.20% more burn events over 30 days, while the net mint-minus-burn event balance rises +1.82 pp. The provider-day version points the same way: total future mint-or-burn origin counts rise +59.11%. Net of the vehicle’s own origin-day V3 action and provider-day counts, the provider-day estimate remains +4.81% at 30 days and +6.02% at 120 days. Splitting the same design by issuer shows that the churn varies by issuer: total V3 LP actions rise by +86.61% for DAI, +41.96% for USDC, and +99.89% for USDT. For USDT, the action response comes from both sides of the book, with future mint events up +98.64% and burn events up +100.25%, and the provider-day response is +91.94%, while the net event balance is negative at -0.66 pp; DAI and USDC have positive net balances, +5.27 pp and +2.38 pp. The V4 modify-liquidity archive sharpens the protocol boundary. Among 1,225,274 five-vehicle V4 vehicle-action assignments, stablecoins account for 56.6%, compared with 31.4% in the V3 mint/burn assignment panel; USDT alone is 22.6% in V4, compared with 11.3% in V3. V4 actions are range-management intensive: 69.3% of assignments are narrow or medium tick ranges, and 4.8% are full-range. Those quantities count events and provider-days; dollar-valued provider flows and unique LP identities remain separate objects, and the V4 comparison maps native ETH to the WETH vehicle family. Within vehicle and origin-date effects, a one log-point higher capital stock is associated with +0.90 pp higher intermediary episode share and +0.50 higher excess-use ratio. The Same pool, Rent incidence, V3 churn, and V3 activity-controlled rows of Table 13 isolate the corresponding capital, yield, event, and provider-day margins. Together, the predictive associations organize capital stocks, venue reach, LP action counts, and vehicle roles around inventory rebalancing, range-management churn, and venue-footprint selection, consistent with constrained liquidity provision in currency markets (Huang et al., 2025). Causal provider-flow effects lie outside the design.

The forward V4 exercise uses the singleton’s own swap and modify-liquidity logs. Within vehicle and origin-date effects, a 10 pp stable route-minus-capital gap predicts +5.40 pp higher multi-leg V4 transaction share, +4.50 pp higher same-asset internal-leg share, and +2.32 pp higher gross-to-net singleton reduction. The corresponding LP-flow specification values each vehicle token side of V4 modify-liquidity events. On V4-active origin days, net of current vehicle-token side flow activity, the same gap predicts +30.91% more future admissible LP flow after 120 days. Both sides of the book move: add flow rises +23.99%, while remove flow rises +40.07%. Net of the vehicle’s current V4 action, origin-count, and sender-count activity, the same gap predicts +30.04% more future V4 modify-liquidity actions and +12.64% more future sender-days. The V4 accounting, V4 LP flow, and V4 activity-controlled rows of Table 13 report these transaction-share, flow, action, and sender-day coefficients. This is the liquidity-provision result connected to the V4 accounting design: the

economic route remains visible even when the singleton nets settlement, and LPs respond on flow, action-count, and sender-day margins (Lehar and Parlour, 2025). The measure is filtered vehicle-token-side flow; whole-pool TVL, LP inventory, and causal provider responses remain outside its scope.

Table 13 consolidates the route-capital-gap regressions. A compact notation for the vehicle-day rows is

$$Y_{c,t+h} = \beta G_{c,t}^S + X'_{c,t} \Gamma + \alpha_c + \delta_t + \varepsilon_{c,t+h}, \quad (15)$$

In Equation (15), $G_{c,t}^S$ is the stablecoin route-minus-capital gap at the origin date. The table reports the association for a 10 pp larger gap. Positive capital-share convergence and incumbent-pool coefficients locate the stock response. V3 activity covaries with mint-and-burn churn, and V4 activity covaries with flash-accounting intensity, vehicle-token-side LP flow, actions, and sender-days, the provider-side margins closest to the constrained-liquidity adjustment in Huang et al. (2025).

Table 14 reports the corresponding route-to-provider mechanism in the opposite direction, using the singleton’s own accounting intensity at the vehicle-day level:

$$Y_{c,t+h} = \beta M_{c,t} + X'_{c,t} \Gamma + \alpha_c + \delta_t + \varepsilon_{c,t+h}. \quad (16)$$

In Equation (16), $M_{c,t}$ is one of the three flash-accounting proxies, $X_{c,t}$ contains origin-day swap activity, current vehicle-token side LP flow, current vehicle-linked TVL, current LP actions, and the current narrow-or-medium range share, and the fixed effects compare vehicles on the same day. Table 14 uses the 120-day horizon because that is the interval over which the earlier V2 and V3 provider responses are most visible. Its Internal same-asset share and Multi-leg transaction share rows forecast more future LP flow, more vehicle-linked TVL, and more LP actions in the first three outcome columns. The range evidence separates value from management churn. Across the two flow-share columns, all three flash proxies predict higher narrow-or-medium shares and lower broad shares. Across the two action-share columns, internal and multi-leg routing predict fewer narrow-or-medium shares and more wide-or-very-wide shares. V4 flash accounting is therefore associated with concentrated value placement and broader repositioning activity, distinct from a one-margin inventory increase.

Table 15 joins the two margins. The row variable is the interaction between the stable route-minus-capital gap and the same-day V4 accounting proxy, with the same vehicle and date effects and the same origin-day activity controls as Table 14. In its first two rows, the Vehicle-linked TVL and LP actions columns rise while the range-share columns move from narrow-or-medium toward wide-or-very-wide placements. The Gross-to-net reduction row reverses those four signs. The result says that the economically visible V4 route margin is connected to provider reallocation, because the accounting proxies have different liquidity meanings. That distinction fits a view of liquidity providers as active inventory and range managers facing fee, adverse-selection, and adjustment-cost tradeoffs (Lehar and Parlour, 2025; Huang et al., 2025).

Table 13: Liquidity-provision outcomes and stable route-capital gaps

Margin	Outcome ($Y_{c,t+h}$)	Horizon	Effect of +10 pp in $G_{c,t}^S$	Obs.	Clusters
V2 stock	Capital share change [pp]	120d	+1.4*** (0.2)	10,595	2,119
V2 stock	Log deposited capital [log points]	120d	+0.205*** (0.031)	10,595	2,119
Portfolio rank	Capital-rank catch-up [rank positions]	120d	+0.225*** (0.036)	10,595	2,119
Same pool	Same-pool capital [log points]	120d	+0.007 (0.009)	174,712	2,113
Pool footprint	Incumbent-pool capital [log points]	120d	+0.124** (0.054)	10,417	2,117
Rent incidence	V3 fee yield [log points]	120d	-0.003 (0.005)	208,360	189
V3 churn	Mint events [log points]	30d	+0.679*** (0.037)	11,195	2,239
V3 churn	Burn events [log points]	30d	+0.642*** (0.035)	11,195	2,239
V3 churn	Net mint-burn balance [events]	30d	+0.018*** (0.002)	11,195	2,239
V3 activity-controlled	Provider-day activity [log points]	120d	+0.060*** (0.010)	11,195	2,239
V4 accounting	Multi-leg transaction share [pp]	same day	+5.4*** (0.6)	2,599	522
V4 accounting	Internal same-asset share [pp]	same day	+4.5*** (0.4)	2,599	522
V4 accounting	Gross-to-net reduction share [pp]	same day	+2.3*** (0.4)	2,599	522
V4 LP flow	Gross vehicle-side flow [log points]	120d	+0.309*** (0.041)	2,592	521
V4 LP flow	Add flow [log points]	120d	+0.240*** (0.036)	2,592	521
V4 LP flow	Remove flow [log points]	120d	+0.401*** (0.047)	2,592	521
V4 activity-controlled	LP actions [log points]	120d	+0.300*** (0.024)	2,599	522
V4 activity-controlled	Sender-days [log points]	120d	+0.126*** (0.016)	2,599	522

Note: Each cell reports β from Equation (15), the association for a 10 pp larger stablecoin route-minus-capital gap $G_{c,t}^S$, with clustered standard errors in parentheses. Capital-share change is target-date minus origin-date share within the five-vehicle V2 capital total. Because origin capital share enters both this outcome and $G_{c,t}^S$, capital-share and capital-rank rows retain a mechanical mean-reversion component and describe convergence; identifying provider reallocation requires exogenous variation. Log-point outcomes are changes in $\log(1+q)$, where q is deposited capital in dollars, event count, provider-day count, sender-day count, vehicle-token-side LP flow in dollars, or fee yield as named by the row. Capital-rank catch-up is origin rank minus future rank, so a positive value is an improvement; net mint-burn balance is the future mint-event count minus burn-event count. V4 accounting outcomes are same-day transaction shares in pp. Vehicle-day rows absorb vehicle and origin-date effects unless the row label implies a pool-level design: same-pool capital absorbs pool-vehicle and origin-date effects, V3 fee yield absorbs pool and origin-date effects, and V4 rows use V4-active origin vehicle-days. V3 activity-controlled and V4 activity-controlled rows include current origin-day action, origin-count, and sender-count controls; V4 flow rows include current vehicle-token-side LP-flow controls. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively. These predictive mechanism regressions use observational variation.

Table 14: V4 flash accounting and subsequent liquidity provision

Flash-accounting proxy ($M_{c,t}$)	Future LP flow [log points]	Vehicle-linked TVL [log points]	LP actions [log points]	Flow narrow/medium [pp]	Flow broad [pp]	Action narrow/medium [pp]	Action wide/very-wide [pp]
Internal same-asset share	+0.035* (0.020)	+0.265*** (0.031)	+0.071*** (0.013)	+1.398*** (0.137)	-1.398*** (0.137)	-2.460*** (0.384)	+2.889*** (0.354)
Multi-leg transaction share	+0.076** (0.032)	+0.123*** (0.018)	+0.058*** (0.017)	+1.165*** (0.104)	-1.165*** (0.104)	-1.867*** (0.203)	+1.808*** (0.224)
Gross-to-net reduction share	+0.139*** (0.041)	+0.026 (0.031)	+0.044** (0.019)	+1.223*** (0.200)	-1.223*** (0.200)	+0.362 (0.373)	+0.021 (0.328)
Observations / date clusters							2,015 / 403
Asset and date effects							Yes
Origin-day activity controls							Yes

Note: Each cell reports the effect of a 10 pp increase in flash-accounting proxy $M_{c,t}$ on the column outcome, with standard errors in parentheses. Multi-leg share is the fraction of vehicle-touching transactions with at least two swap legs; internal same-asset share is the fraction in which the vehicle asset appears on at least two legs; gross-to-net reduction is $1 - |\sum_{\ell} \Delta q_{c\ell}| / \sum_{\ell} |\Delta q_{c\ell}|$, aggregated from vehicle-token amount changes within each singleton transaction. The unit is a vehicle-day in Uniswap v4. All columns use the 120-day horizon, absorb vehicle and origin-date fixed effects, include origin-day swap activity, current vehicle-token side LP flow, current vehicle-linked TVL, current LP actions, and current range-composition controls, and cluster by origin date. LP flow is the admissible dollar value of vehicle-token sides in modify-liquidity events. Vehicle-linked TVL is full-pool provider TVL linked to a vehicle token side after address-resolution and physical TVL/volume bounds. Flow-share columns partition future dollar-valued modify-liquidity flow into narrow-or-medium versus broad ranges; action-share columns partition future modify-liquidity events into narrow-or-medium and wide-or-very-wide placements. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 15: Stable shortfalls, V4 flash accounting, and LP repositioning

Interaction term ($G_{c,t}^S \times M_{c,t}$)	Vehicle-linked TVL [log points]	LP actions [log points]	Narrow/medium ranges [pp]	Wide/very-wide ranges [pp]
Stable gap $\times$ internal same-asset share	+0.292*** (0.079)	+0.082*** (0.027)	-6.531*** (0.720)	+6.189*** (0.657)
Stable gap $\times$ multi-leg transaction share	+0.225*** (0.035)	+0.132*** (0.013)	-4.539*** (0.442)	+4.438*** (0.381)
Stable gap $\times$ gross-to-net reduction share	-0.667*** (0.105)	-0.121*** (0.043)	+4.411*** (1.261)	-4.405*** (1.126)
Observations / date clusters				1,209 / 403
Asset and date effects				Yes
Origin-day activity controls				Yes

Note: Each cell reports $G_{c,t}^S \times M_{c,t}$, the interaction of a 10 pp larger stable route-minus-capital gap with a 10 pp larger V4 accounting proxy; $G_{c,t}^S$ is stablecoin route share minus deposited-capital share and the three $M_{c,t}$ proxies are defined in Table 14. Log outcomes are in log points; range outcomes are in pp. The unit is a stable vehicle-day in Uniswap v4. All columns use the 120-day horizon, absorb vehicle and origin-date fixed effects, include origin-day swap activity, current vehicle-token side LP flow, current vehicle-linked TVL, current LP actions, and current range-composition controls, and cluster by origin date. Vehicle-linked TVL is full-pool provider TVL linked to a vehicle token side after address-resolution and physical TVL/volume bounds. Range shares partition future modify-liquidity events into narrow-or-medium, wide-or-very-wide, and full-range placements. The estimates are predictive associations from observational variation. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 16: Stable shortfalls, V4 flash accounting, and LP-flow sides

Interaction term ($G_{c,t}^S \times M_{c,t}$)	Gross LP flow [log points]	Add-side LP flow [log points]	Remove-side LP flow [log points]
Stable gap $\times$ internal same-asset share	-0.327*** (0.057)	-0.303*** (0.057)	-0.354*** (0.061)
Stable gap $\times$ multi-leg transaction share	-0.073** (0.035)	-0.097*** (0.034)	-0.048 (0.037)
Stable gap $\times$ gross-to-net reduction share	+0.214** (0.090)	+0.156* (0.090)	+0.276*** (0.095)
Observations / date clusters			1,209 / 403
Asset and date effects			Yes
Origin-day activity controls			Yes

Note: Each cell reports $G_{c,t}^S \times M_{c,t}$, the interaction of a 10 pp larger stable route-minus-capital gap with a 10 pp larger V4 accounting proxy; $G_{c,t}^S$ is stablecoin route share minus deposited-capital share and the three $M_{c,t}$ proxies are defined in Table 14. Outcomes are future log admissible vehicle-token side modify-liquidity flow over 120 days. The unit, fixed effects, controls, covariance, and stablecoin sample match Table 15. Add and remove flows value the vehicle-token sides that can be assigned through the V4 modify-liquidity records. These flow-side predictive specifications measure vehicle-token flows; whole-pool inventory and causal LP supply remain outside their scope. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively.

The flow side identifies repositioning separately from provider-token flow. Table 16 applies the same interaction design to future vehicle-token side modify-liquidity flow. Its Internal same-asset share row is negative in the gross, add-side, and remove-side columns; the Multi-leg transaction share row is negative for gross and add-side flow and smaller for remove-side flow. The Gross-to-net reduction row has positive coefficients in all three columns, with the largest estimate on remove-side flow. Read beside Table 15, these patterns locate the response in stock, action, and range-management margins, while net settlement compression is closer to turnover in the vehicle-token side of modify-liquidity events.

Table 17 makes the protocol comparison directly on the same vehicle and day. It restricts the sample to origins with observed V4 singleton swap activity and stacks the V3 mint/burn outcome and the V4 modify-liquidity outcome. For protocol $p \in \{V3, V4\}$, vehicle v , origin date t , and horizon h , the conditional contrast is

$$Y_{pvt+h} = \alpha_{v,t} + \delta_p + \beta_h \mathbf{1}\{p = V4\} \mathbf{1}\{v \in \mathcal{S}\} G_{v,t} + \gamma'_h \mathbf{X}_{pvt} + \varepsilon_{pvt+h}, \quad (17)$$

In Equation (17), $G_{v,t}$ is route share minus deposited-capital share, $\mathcal{S}$ is the stablecoin set, and $\mathbf{X}_{pvt}$ contains the same protocol's current LP actions and active origins. Vehicle-date effects compare V3 and V4 for the same vehicle and day, protocol effects absorb level differences, and standard errors cluster by origin date. In the 120-day column of Table 17, a 10 pp stable route-minus-capital shortfall is associated with a V4-minus-V3 LP-action differential of 0.320 log points and an active-origin differential of 0.115 log points. The comparison estimates conditional protocol differences; the causal effect of V4 adoption remains outside the design. Its narrower role is to distinguish the interaction from a higher level of counted V4 activity: conditional on the same vehicle-date origin, the stable shortfall is more strongly linked to future provider action in V4 than in V3.

The dollar-flow analogue disciplines that reading. Table 18 values the vehicle-token side of

Table 17: Stable shortfalls and the V4-minus-V3 LP-action response

Outcome [log points]	7 days	30 days	120 days
LP actions	+0.163*** (0.016)	+0.212*** (0.015)	+0.320*** (0.013)
Active origins	+0.045*** (0.011)	+0.062*** (0.011)	+0.115*** (0.010)
Stacked observations / dates			5,198 / 522
Asset-date and protocol effects			Yes
Current protocol LP controls			Yes

Note: Each cell reports the V4-minus-V3 difference in the association between a 10 pp larger stable route-minus-capital gap and the future LP-action outcome, with standard errors in parentheses. Outcomes are log points. The stacked unit is a protocol–vehicle–origin–day row, restricted to vehicle-origin days with observed V4 singleton swap activity. All columns absorb vehicle-date and protocol fixed effects, include same-protocol current LP-action and active-origin controls, and cluster by origin date. V3 outcomes are mint/burn event counts and active origins. V4 outcomes are modify-liquidity event counts and active origins. The estimates are conditional protocol differences and do not identify the effect of V4 adoption. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively.

V3 mint/burn events and V4 modify-liquidity events with the same daily price panel, stacks the two protocols on the same V4-active vehicle dates, and repeats the vehicle-date comparison for future LP flow. The protocol difference is concentrated in delayed removal turnover and short-run range reallocation. In the Remove-side flow row, the 120-day differential is 0.046 log points for a 10 pp stable shortfall. The Narrow/medium flow-share row rises by 1.5 pp at 7 days and 1.1 pp at 30 days, with equal broad-share reductions and a smaller same-sign 120-day effect. The same shortfalls are therefore associated with more V4 provider activity and, on the dollar side, with delayed turnover and range reallocation. This is why the mechanism is described as range management and reallocation, consistent with concentrated-liquidity provider adjustment costs in [Lehar and Parlour \(2025\)](#), distinct from a measured capital-supply shock.

The stock-and-footprint analogue changes the economic reading. Table 19 measures vehicle-linked reported full-pool TVL and pool count in both protocols on the same V4-active vehicle dates. In its Reported TVL growth column, a 10 pp stable route-minus-capital shortfall is associated with V4-minus-V3 differentials of 0.237, 0.297, and 0.637 log points after 7, 30, and 120 days. The Pool-footprint growth column rises from 0.075 to 0.149 and 0.275 log points over the same rows. These monotone stock and footprint responses sit beside flat gross and add-side vehicle-token flows, more provider actions, delayed removal turnover, and short-run range reallocation. Taken together, the results associate the V4 environment with provider reorganization across pools and positions, alongside gradual vehicle-token-side deposit adjustment. The quantity remains bounded: a pool contributes its full reported TVL to every vehicle it contains, so vehicle–vehicle pools can enter both vehicle series. The measure is reported full-pool TVL; side-specific LP inventory and reconstructed deposited capital remain separate quantities. The protocol contrast identifies conditional differences; adoption effects remain outside the design. Within those limits, the persistent footprint response is economically substantial and connects stable vehicle shortfalls to an LP allocation margin consistent with the endogenous pool adjustment in [Lehar and Parlour \(2025\)](#).

Table 18: Stable shortfalls and the V4-minus-V3 LP-flow response

Outcome	7 days	30 days	120 days
Gross LP flow [log points]	−0.043 (0.033)	−0.009 (0.027)	+0.029 (0.022)
Add-side flow [log points]	−0.043 (0.033)	−0.023 (0.027)	+0.012 (0.023)
Remove-side flow [log points]	−0.032 (0.034)	+0.007 (0.028)	+0.046** (0.022)
Narrow/medium flow share [pp]	+1.530*** (0.236)	+1.126*** (0.133)	+0.329*** (0.095)
Broad flow share [pp]	−1.530*** (0.236)	−1.126*** (0.133)	−0.329*** (0.095)
Stacked observations / dates	5,188 / 521		
Asset-date and protocol effects	Yes		
Current protocol LP-flow controls	Yes		

Note: Each cell reports the V4-minus-V3 difference in the association between a 10 pp larger stable route-minus-capital gap and future vehicle-token-side LP flow, with standard errors in parentheses. Flow-amount outcomes are log points; range-composition outcomes are pp. The stacked unit is a protocol–vehicle–origin–day row, restricted to vehicle-origin days with observed V4 singleton swap activity. All columns absorb vehicle-date and protocol fixed effects, include same-protocol current gross, add, remove, and sender-day flow controls, and cluster by origin date. V3 outcomes value mint and burn event token sides. V4 outcomes value modify-liquidity event token sides. The estimates are conditional protocol differences in vehicle-token flows and do not identify the effect of V4 adoption. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively.

Table 19: Stable shortfalls and the V4-minus-V3 reported-liquidity footprint

Horizon	Reported TVL growth [log points]	Pool-footprint growth [log points]
7 days	+0.237*** (0.073)	+0.075*** (0.017)
30 days	+0.297*** (0.056)	+0.149*** (0.020)
120 days	+0.637*** (0.042)	+0.275*** (0.022)
Obs. / date clusters (7d)	2,986 / 301	
Obs. / date clusters (30d)	2,756 / 278	
Obs. / date clusters (120d)	1,856 / 188	
Asset-date and protocol effects	Yes	
Origin TVL and pool-count controls	Yes	

Note: Each cell reports the V4-minus-V3 difference in the association between a 10 pp larger stable route-minus-capital gap and future growth in vehicle-linked reported full-pool TVL or pool count, with standard errors in parentheses. Outcomes are changes in log one plus the reported quantity. The stacked unit is a protocol–vehicle–origin–day row, restricted to V4-active vehicle dates with strict origin and target support in both protocols. All columns absorb vehicle-date and protocol fixed effects, include same-protocol origin TVL and pool-count controls, and cluster by origin date. A pool contributes its full reported TVL when it contains the vehicle; vehicle–vehicle pools can enter both vehicle series. The quantity is a reported stock-and-footprint proxy; reconstructed deposited capital, side-specific LP inventory, and causal protocol effects remain outside its scope. Asterisks *, **, and *** denote statistical significance at the 10%, 5%, and 1% levels, respectively.

5.5 Protocol architecture

Protocol design changes the set of routes that can be formed. Uniswap v1 created pools between an ERC-20 token and ETH, so token-to-token trades crossed ETH; Uniswap v2 allowed arbitrary ERC-20–ERC-20 pools and made direct token-to-token trading possible.⁶ Forced intermediation was material while the constraint lasted: token-to-token trades routed through ETH account for 8.0% of Uniswap v1 swap transactions and 8.1% of its ETH volume over the venue’s life, and roughly one trade in ten in 2019 and 2020, when the venue was most active. A transaction from the crash weekend of March 2020 shows the mandate operating: with no direct pool between MKR and DAI, a trader selling 250 MKR received 439.13 ETH at one exchange contract and spent the identical ETH amount at a second for 49,892.40 DAI.⁷ Native-asset pairing remained pervasive after this mandate ended. The share of single-leg Uniswap v2 trades executed in WETH pools was 95.1% in 2020 and 95.5% in 2026, remaining between 93.4% and 97.9% in every intervening year. Among the 477,633 atomic pairs that ever traded on Uniswap v2, 97.1% include WETH, and WETH’s share of newly traded atomic pairs rose from 84.1% in 2020 to 97.9% in 2026. This persistence is consistent with deep WETH pools, launch conventions, and inherited search paths. Separating those forces from the original protocol constraint requires comparison among atomic pairs that had the same alternatives available.

Uniswap v4 separates route choice from physical settlement more sharply. All pools reside in one singleton, and flash accounting records balance changes internally before settling the final net amounts at the end of a call.⁸ An economic route can therefore cross an intermediary pool without generating an external token transfer at every step. Decoded pool actions preserve the trading path despite this netting, while the architecture changes how liquidity, fees, and settlement are organised around that path.

The route evidence leaves an explanation six features to account for. The stablecoin gain occurs mainly across ultimate pairs, with little repeated substitution within matched ultimate pairs; it appears on short and long paths alike and on routes that stay within one exchange as well as routes that do not; the incidence of intermediation remained flat when the software that searches for paths was rewritten, though the exchange set it reached did widen; the gain is present where every leg was priced by the same unchanged constant-product rule, and absent from the venue whose invariant is built for pegged assets, which carries no intermediation at all; local two-leg depth predicts which vehicle carries route mass inside an opportunity; and the measured WETH-pairing persistence shows that route opportunities can outlast the protocol rule that first created them. Together they narrow the set of viable explanations. A viable explanation must instead account for why newly active and growing ultimate pairs attach themselves to stablecoins.

⁶See Adams, Zinsmeister, and Robinson, *Uniswap v2 Core*, pp. 1–2.

⁷Transaction [0x4dca160d...96a0ca16](#), block 9,674,728, March 15, 2020. The trace is selected by a deterministic rule, the largest ETH amount routed among the 129,810 mandate-era token-to-token transactions whose two ETH legs match exactly, and the token identities are verified against the public transfer record.

⁸See Adams et al., *Uniswap v4 Core*, Sections 1 and 3.

6 Conclusion

Vehicle dominance on decentralised exchanges has shifted markedly from native assets toward stablecoins through a reorganisation of the trading network. Across matched ultimate pairs without a route-feasibility condition, positive and negative stable-for-native movements offset. Aggregate dominance rises as activity reallocates toward ultimate pairs that already route through stablecoins and as ultimate pairs enter and leave the comparison. The bridge-establishment sample describes the local liquidity margin. During the first month, stable weak-leg depth is below one-tenth of WETH depth on 7,752 of 11,327 comparable ultimate-pair days. Within a bridge event, a 1 pp increase in stable relative depth predicts +0.64 pp higher stable route share. That share is 53.0% when stable depth is at least WETH depth and 69.9% when stable depth is at least twice WETH depth. Native assets regain share in 2023 and 2024 before stablecoins recover the value lead, and USDT’s intermediary use rises relative to its appearance at the observed endpoints of pool routes.

Vehicle-currency theories usually emphasise a reinforcing choice for a given ultimate pair of assets: trading deepens an intermediary’s markets, lower costs attract more trading, and the vehicle role persists. The decomposition adds the network distribution of trading to this account. A vehicle gains aggregate share when it becomes more common within an ultimate pair, when activity shifts toward ultimate pairs that already use it, or when newly active ultimate pairs route through it. In our sample, the latter two margins dominate. This distinction changes how aggregate vehicle shares should be interpreted: a rising share can reflect a reorganisation of trading toward parts of the network in which the challenger already connects markets even when positive and negative changes inside continuing ultimate pairs cancel. Models and empirical tests that hold the ultimate pair fixed describe the first margin; aggregate turnover combines both.

The mechanism evidence locates where such a role is made. Stable vehicles appear in new and repeatedly active ultimate pairs, and their route share rises with the relative deposited depth of the two-leg bridge. Deposited capital is a proximate liquidity-formation state variable; executable depth and provider-supply shocks are distinct objects. The entry analysis separately documents persistence in the first intermediary’s identity. Liquidity provision moves on a slower margin: deposited capital remains native-heavy, route-capital gaps covary with later capital-share convergence and provider activity, and the USDC/SVB window shows stable-route demand rising while capital shifts toward WETH. The capital-share convergence regressions retain a mechanical mean-reversion component and estimate the convergence association. Route opportunities, local depth, entry identity, and provider activity therefore move together on distinct horizons.

The accounting extends beyond decentralised exchanges. A currency’s share of settled trade can decline while every trading relationship that uses it today continues to do so, provided trade grows fastest among relationships that never adopted it. The renminbi arrangements between China and Brazil described in Section 1 illustrate how a new bilateral settlement link can change the set of active ultimate pairs. Aggregate currency shares are always measured over a set of ultimate pairs that is itself changing, and the two margins quantified here are the margins through which such a decline would run. Records that retain only atomic legs confound these composition margins with

a change of intermediary within an ultimate pair; a transaction log that preserves the whole path separates them.

The economic importance of a vehicle comes from the trading it connects. Vehicle dominance is therefore jointly determined by the intermediary selected for each ultimate pair and by the distribution of routed activity across the network.

A From contract events to economic routes

Route construction begins with the contract and event format of each exchange. We identify pools from the protocol’s factory or registry, map contract addresses to tokens, and convert raw integer amounts using the token precision that applied at the time of the trade. Contract addresses define token identity because unrelated tokens can share a symbol and the same economic asset can appear in native and wrapped form. ETH and WETH are consolidated only after this address-level reconstruction; every other token remains distinct unless a stated economic classification combines it. Within each transaction, atomic trades (pool-swap legs) are ordered by block, transaction, and log position. The resulting token flows reveal the exchange that occurred. A single connected chain gives one route and therefore one ultimate trade between its endpoints; parallel chains identify an order split, and disconnected event groups remain separate. Because one transaction can contain unrelated calls as well as trades across several exchanges, this procedure also prevents two simultaneous routes from being mistaken for one longer route. The output records the sequence of atomic pairs and pools traversed, together with the ultimate pair. It does not infer that sequence from the transaction sender or router label.

A complete token sequence is sufficient for the route-count sample. Value weighting additionally requires coherent amounts on every leg and a common dollar valuation. A failed valuation removes the observation from the value calculation while preserving the observed route. Appendix F gives the separate exclusion for self-returning paths.

Historical token precision cannot be verified for 13 tokens traded in 22 constant-product pools. Reserve-reconstruction analyses omit those pools. An upper-bound check removes every route involving any of the 13 tokens: the resulting bound is 5,950 routes and \$10.91 million of route value. Only 1,036 of those routes use one of the affected tokens as an intermediary, representing \$1.284 million, or 0.00043% of value on routes whose legs can all be valued. All 13 tokens lie outside the native and stablecoin groups used in the main comparison. The exclusion affects exact reserve reconstruction; the route-composition result is unchanged.

B Pricing models and validation

The counterfactual comparison requires an exact-input quote from each pool’s state and the proposed input amount. Because pool families use different invariants and arithmetic, we validate their pricing models separately. For selected executed swaps, we reconstruct the state immediately beforehand and compare predicted with realised output. Curve’s amplification coefficient and Balancer’s weight ratio are unavailable in the historical subgraph records, so we estimate them on one set of trades and evaluate the resulting quotes on held-out trades. A route enters the cost comparison only when the relevant pool states and pricing models are available before the transaction.

B.1 Constant-product pools

Uniswap v2 and SushiSwap v2 price the product of reserves (Lehar and Parlour, 2025). Let R_{i,p,e^-} and R_{o,p,e^-} denote pool p 's input- and output-token reserves immediately before transaction e , Δ_i the token input, f_p the fee rate, and $\varphi_p = 1 - f_p$ the fraction of input retained for pricing. An exact-input swap moves the pool along

$$(R_{i,p,e^-} + \varphi_p \Delta_i)(R_{o,p,e^-} - \Delta_o) = R_{i,p,e^-} R_{o,p,e^-},$$

which inverts in closed form to

$$\Delta_o = \frac{\varphi_p \Delta_i R_{o,p,e^-}}{R_{i,p,e^-} + \varphi_p \Delta_i}, \quad \varphi_p = \frac{997}{1000},$$

the 30 basis point fee both venues charge, applied on chain as the integer ratio above.⁹

A multi-leg route composes these one hop at a time, threading the output token of one pool into the next. The cost comparison evaluates the dollar outputs of the best direct and intermediated routes at the same state and dollar input notional,

$$10^4 \Delta C_{i,o,k,q,e}^D = 10^4 \frac{O_{i,o,q,e}^D - O_{i,o,k,q,e}^I}{O_{i,o,q,e}^D}.$$

Both outputs must be computed from one reconstructed pre-transaction state. A realised trade supplies the output of the chosen route alone; the direct and intermediary alternatives require reconstructed states for every pool in the comparison. Quoting each executed constant-product swap against the state immediately before it gives a median absolute error of 0.0000% across 8,024 swaps; 96.7% of quotes are within 1% and 95.2% are within 0.01% of realised output.

B.2 Concentrated liquidity

Uniswap v3 distributes liquidity across price ranges, so executable depth depends on the active range and on the initialized ticks reached by a trade (Adams et al., 2021). Within one active range, let $\tilde{R}_{i,p,e^-}$ and $\tilde{R}_{o,p,e^-}$ denote the virtual reserves of the input and output assets. Liquidity L_{p,e^-} and the output-per-input marginal price ρ_{p,e^-} satisfy

$$\tilde{R}_{i,p,e^-} = \frac{L_{p,e^-}}{\sqrt{\rho_{p,e^-}}}, \quad \tilde{R}_{o,p,e^-} = L_{p,e^-} \sqrt{\rho_{p,e^-}}.$$

If f_p is the pool fee and $\varphi_p = 1 - f_p$, an exact input Δ_i returns

$$\Delta_o = \frac{\varphi_p \Delta_i \tilde{R}_{o,p,e^-}}{\tilde{R}_{i,p,e^-} + \varphi_p \Delta_i}$$

⁹The SushiSwap v2 atomic-pair contract implements the same adjustment; see [UniswapV2Pair.sol](#).

Table 20: Concentrated-liquidity quoter, by direction and tick crossing

Validation group	Swaps	Median	90th pct.	Maximum	Within 1%
token0 in, no tick crossed	2,290	0	0	9.98×10^{-4}	100.0%
token0 in, tick crossed	57	0	0	3.45×10^{-7}	100.0%
token1 in, no tick crossed	1,816	0	1.69×10^{-8}	1.48×10^{-4}	100.0%
token1 in, tick crossed	55	0	5.44×10^{-9}	8.74×10^{-9}	100.0%
Pooled	4,218	0	1.15×10^{-14}	9.98×10^{-4}	100.0%

Note: Absolute quote error in percent of realised output. The sample contains 4,218 realised Uniswap v4 swaps on three dates across 13 pools and nine atomic pairs, at static fee tiers 0, 7, 8, 100, 500, 10,000 and 20,000. Each quote uses liquidity events ordered strictly before the swap and the price state left by the preceding swap in the same pool.

until the trade reaches the next initialized tick. Range boundaries lie on a geometric grid, and crossing tick j updates active liquidity by its signed change,

$$\sqrt{\rho(j)} = 1.0001^{j/2}, \quad L_{p,e^-} \mapsto L_{p,e^-} \pm \Delta L_j,$$

with the sign determined by the direction of the trade. The calculation then continues from the new active range until the input is exhausted. The quoter uses the contract’s integer Q96 arithmetic, rounding input amounts up and output amounts down.

The active tick map is accumulated from every mint and burn strictly before the validation day, after which same-day liquidity changes and swaps are replayed together by block and log index. Consecutive swaps in one pool supply the price state: the earlier event leaves `sqrtPriceX96` and `tick` immediately before the later swap, while the interleaved event stream supplies active liquidity. We report errors separately by input direction and tick crossing because those cases use different arithmetic branches. Table 20 applies this validation to Uniswap v4. V4 retains concentrated-liquidity pricing for static-fee pools but implements it through a singleton and flash accounting.¹⁰ The validated static fee tiers include 0, 7, 8, and 20,000 hundredths of a basis point. Hook-bearing and dynamic-fee pools are excluded because the raw records do not identify their transaction-level cash flows and realised fees; Appendix E reports their share of v4 trading.

Every direction and tick-crossing row of Table 20 reproduces realised output within 0.001% in the Maximum column.

B.3 StableSwap pools

Curve interpolates between a constant-sum curve, which is flat and charges no slippage, and the constant-product curve of Section B.1.¹¹ For n tokens with normalized pre-transaction balances

¹⁰See Adams et al., *Uniswap v4 Core*, Sections 1 and 3.

¹¹See Egorov, *StableSwap—Efficient Mechanism for Stablecoin Liquidity*.

R_{j,p,e^-} , amplification A_p , and invariant D_p ,

$$A_p n^n \sum_j R_{j,p,e^-} + D_p = A_p D_p n^n + \frac{D_p^{n+1}}{n^n \prod_j R_{j,p,e^-}},$$

where $A_p \rightarrow 0$ recovers the constant product and $A_p \rightarrow \infty$ recovers the constant sum. The contract solves D_p by Newton iteration and then solves the same invariant for the post-trade output balance $R'_{o,p,e}$ after R_{i,p,e^-} is increased by Δ_i . Defining $D_\Pi = D_p \prod_j D_p / (n R_{j,p,e^-})$, the two loops are

$$D_p \leftarrow \frac{(A_p n \sum_j R_{j,p,e^-} + n D_\Pi) D_p}{(A_p n - 1) D_p + (n + 1) D_\Pi}, \quad R'_{o,p,e} \leftarrow \frac{(R'_{o,p,e})^2 + c}{2 R'_{o,p,e} + b - D_p},$$

with b and c formed from the balances of the tokens the swap does not touch. Balances are rescaled to eighteen decimals before either loop runs, because the invariant treats its tokens as interchangeable at par and a six-decimal stablecoin read on an eighteen-decimal scale is off by a factor of 10^{12} . Output is net of the pool fee,

$$\Delta_o = (R_{o,p,e^-} - R'_{o,p,e} - 1) \varphi_p,$$

carrying the contract's own off-by-one guard on the raw difference.

A_p is Curve-specific and the standardised subgraph schema drops it, so it is identified from data instead of fetched. Every other quantity in the invariant is observed and every Curve swap reports both amounts, which leaves A_p as the only unknown,

$$\hat{A}_p = \arg \min_{A_p} Q_{0.90} \left\{ \frac{|Q_p(i \rightarrow o, \Delta_i; \chi_{p,e^-}(A_p)) - \Delta_o|}{\Delta_o} \right\},$$

fitted on the first half of a pool-day's trades over a logarithmic grid with a ternary refinement, then evaluated on the trades not used for estimation. A pool-day is retained when the 90th percentile of fitting error is below 1%. Using the upper tail prevents a partial fit from passing: one misspecified pool produced a median fitting error below 0.1% but a 34% median error on the other trades because the invariant fit only part of its activity. Curve has also operated crypto-pools since 2021 that follow the CryptoSwap invariant. Applying StableSwap to these pools yields a best-fit A_p of 3 and a 36% median quote error, so pool-days whose observed trades are inconsistent with StableSwap remain outside this component exercise.

Table 21 reports the held-out fit by validation day.

Across four validation days, held-out median errors range from 0.029% to 0.039%.

Table 21: StableSwap calibration and held-out quote error, by validation day

Day	Pools fitted	Pools excluded	Held-out trades	Median $\hat{A}_p$	Median error	Within 1%
2021-05-28	9	14	367	12,716	0.030%	100.0%
2022-09-12	11	16	262	20,493	0.029%	100.0%
2023-12-28	20	28	471	23,428	0.039%	98.9%
2025-04-13	42	50	1,207	46,851	0.036%	99.7%

Note: The sample contains four sampled days. $\hat{A}_p$ is reported in the module’s convention, which carries the on-chain precision factor of 100. Median error is the median absolute deviation of the quote from realised output on trades the calibration never saw.

B.4 Weighted pools

Balancer weighted pools use a weighted geometric-mean invariant,¹²

$$K_{w,p} = \prod_j R_{j,p,e^-}^{w_{j,p}}, \quad \sum_j w_{j,p} = 1,$$

which gives the two-token exact-input form

$$\Delta_o = R_{o,p,e^-} \left[1 - \left(\frac{R_{i,p,e^-}}{R_{i,p,e^-} + \varphi_p \Delta_i} \right)^{w_{i,p}/w_{o,p}} \right].$$

Only the ratio $w_{i,p}/w_{o,p}$ enters. An equal-weight Balancer pool and a constant-product atomic pair with the same reserves therefore return the same quote. Balancer’s weighted-pool contract accepts inputs up to 30% of the input-token reserve; that contract limit governs whether a quote can execute, but it is not the empirical support rule.¹³ The counterfactual design applies the common 5% own-price-impact cutoff in Equation (11) to Balancer and every other pool family. The subgraph reports the weight ratio $\theta_p = w_{i,p}/w_{o,p}$ and the swap fee at its head block while balances come from a historical snapshot, so a pool that has repriced itself since is quoted with a current parameter against a past state. Liquidity-bootstrapping pools move weights on a schedule by design and managed pools move them on demand, and a weighted pool’s owner can reset the fee at will. What such a pool produces is a small and nearly constant relative bias, which appears directly in the preceding validation: reported parameters reproduce trades in 2024 to a millionth of a basis point and trades in 2022 only to 0.19%. A gap of that size is a fee. The fee enters the quote through $\varphi_p \Delta_i$, so a fee wrong by a basis point moves output by about a basis point, whereas a mistake in the invariant would leave an error that scales with trade size.

We therefore treat both as unknowns to be recovered from the trades a pool actually executed, the move Curve’s amplification coefficient needed in Section B.3, and we keep a reported value only where it survives that test. Reconstructing a pool’s state from its own event history follows [Lehar and Parlour \(2025\)](#); what the weighted family adds is that two arguments of the pricing function

¹²See Martinelli and Mushegian, *A Non-Custodial Portfolio Manager, Liquidity Provider, and Price Sensor*.

¹³See the input-ratio bound in Balancer’s [WeightedMath.sol](#).

are themselves missing from the historical record, so a reconstructed state alone does not deliver a quote. Every quantity in the weighted invariant other than θ_p and f_p is observed, and every Balancer swap reports both amounts, so with balances reconstructed either parameter is the only unknown. We order a pool-day's trades and assign them alternately to a fitting set $\mathcal{O}_p$ and an evaluation set, and every error reported in Table 22 is measured on the evaluation set. For a trial parameter pair ϑ ,

$$\mathcal{E}_p(\vartheta) = Q_{0.90} \left\{ \frac{|Q_p(i \rightarrow o, \Delta_i; \chi_{p,e^-}(\vartheta)) - \Delta_o|}{\Delta_o} : e \in \mathcal{O}_p \right\}$$

is the fitting error, defined only when at least nine tenths of $\mathcal{O}_p$ returns a quote at all and undefined otherwise. The upper tail does the gating for the reason it does at Curve, and the coverage condition stops a parameter that prices one corner of a pool's activity from standing for the pool.

An atomic pair's weight ratio is fitted on the trades crossing it in either direction, with a trade running from b to a scored at the reciprocal exponent, which is what the invariant implies and which doubles the trades available to pin one number,

$$\hat{\theta}_{ab,p} = \arg \min_{\theta} \mathcal{E}_p(\theta \text{ on } \mathcal{O}_p^{a \rightarrow b}, \theta^{-1} \text{ on } \mathcal{O}_p^{b \rightarrow a}),$$

searched over a logarithmic grid spanning Balancer's own weight bounds and refined by ternary search in the log exponent, and attempted only for atomic pairs carrying at least four fitting trades, since one or two are fitted exactly by construction and then read as an exact fit. The fee $\hat{f}_p$ is recovered the same way over the fee range the vault permits. Reading is tried before fitting and the fee before the weights, so each tier adds at most one free scalar and a pool that reproduces its trades on reported parameters is never handed a fitted one,

$$\hat{\vartheta}_p = \begin{cases} (\theta_p, f_p) & \text{if } \mathcal{E}_p(\theta_p, f_p) \leq \bar{e}, \\ (\theta_p, \hat{f}_p) & \text{else if } \mathcal{E}_p(\theta_p, \hat{f}_p) \leq \bar{e}, \\ (\hat{\theta}_{\cdot,p}, f_p) & \text{else if } \mathcal{E}_p(\hat{\theta}_{\cdot,p}, f_p) \leq \bar{e}, \\ \text{excluded} & \text{otherwise,} \end{cases} \quad \bar{e} = 0.1\%. \quad (18)$$

The label a pool carries in the source plays no part in this. A pool-day whose observed trades no tier reproduces is excluded whatever its `poolType` says, and the third tier further requires the fitted ratios to cover nine tenths of the fitting trades, so a pool-day is never accepted on the atomic pairs that happened to fit and then quoted on the rest.

The frequency with which the parameter search binds measures the sufficiency of reported pool state. Of the 256 pool-days accepted across the twelve validation days, 209 clear the gate on reported parameters with nothing identified, 43 need the swap fee alone, and 4 need per-atomic-pair weight ratios. Four fifths of the priced sample therefore carries no fitted parameter, and identification repairs a stale field in a minority of pool-days instead of supplying the invariant.

The gate $\bar{e}$ is set at 0.1% and not at the 1% Curve uses, because the achieved errors separate

Table 22: Weighted-pool quoter and reserve-timing alternatives

Day	Held-out trades	Backward	Flat	Opening	Pools priced / excluded
2021-04-25	4	4.7×10^{-9}	3.00%	3.94%	1 / 0
2021-09-29	273	2.5×10^{-8}	1.24%	0.80%	18 / 8
2022-03-09	524	1.6×10^{-4}	1.05%	2.66%	25 / 4
2022-08-16	701	1.1×10^{-9}	0.75%	1.04%	33 / 12
2023-01-24	832	2.0×10^{-10}	3.73%	4.71%	44 / 11
2023-07-03	301	7.1×10^{-10}	1.78%	3.32%	26 / 19
2023-12-11	185	2.8×10^{-9}	1.77%	2.91%	17 / 20
2024-05-20	265	3.8×10^{-10}	4.81%	5.17%	15 / 36
2024-10-27	142	1.8×10^{-10}	1.68%	2.24%	12 / 36
2025-04-06	248	2.1×10^{-10}	5.32%	9.48%	19 / 40
2025-09-13	501	2.5×10^{-4}	1.09%	1.71%	27 / 25
2026-02-21	582	8.7×10^{-9}	0.68%	1.90%	19 / 6

Note: Median absolute quote error in percent of realised output, on trades no fit ever saw. *Backward* reconstructs per-swap balances by rolling the day’s event sequence back from the closing snapshot, *Flat* keeps the snapshot balances constant across the day, and *Opening* treats the snapshot as the day’s opening state and rolls forward. The sample contains twelve days from 2021-04 to 2026-02.

cleanly there. The threshold separates weighted pools, whose held-out errors range from 10^{-9} to 0.004%, from stable, composable-stable, and elliptic-curve pools, whose held-out errors range from 0.05% to 1.08% and have a 99th percentile of 12%. At a 1% gate the second group clears on a fitted weight ratio and arrives quoted instead of excluded, which is the Curve lesson in a second form: there the over-permissive object was a parameter range and here it is an error threshold.

Across the twelve date rows of Table 22, the Backward column returns a median error of 0.0000% and the Pools priced / excluded column identifies 256 of 473 testable pool-days as weighted pools. The other pool-days carry a median 91.8% of tested swap volume and follow different Balancer pricing rules. The exercise therefore validates the model for identified weighted pools, not for Balancer as a whole.

C Reconstructing pool balances before each trade

Every counterfactual quote uses the pool state immediately before the trade, whereas some historical records store balances only at the end of an hour or day. We recover the earlier state by ordering every balance-changing event within the period and unwinding the sequence from the recorded closing balance. The exercises below validate that procedure on selected constant-product and weighted-pool observations. The cost comparison retains only routes whose pools can be reconstructed in this way.

Let $\mathbf{R}_p^{\text{end}}$ be pool p ’s stored period-end reserve vector and $\Delta_{p,e}$ the vector of net signed token

amounts in event e . The state before event e is the stored value with every later event removed,

$$\mathbf{R}_{p,e}^{\text{pre}} = \mathbf{R}_p^{\text{end}} - \sum_{e' \geq e} \Delta_{p,e'} = \mathbf{R}_{p,e+1}^{\text{pre}} - \Delta_{p,e},$$

which is a single reverse pass over the period. The alternative reading takes the previous period’s stored value as the current period’s opening state and replays forward,

$$\tilde{\mathbf{R}}_{p,e}^{\text{pre}} = \mathbf{R}_{p,\text{prev}}^{\text{end}} + \sum_{e' < e} \Delta_{p,e'},$$

which accumulates every unrecorded balance change between the two stored points. Replaying forward returns a median absolute quote error of 11.8% against realised outputs, whereas unwinding backward returns 0.0000%.

Liquidity additions and withdrawals also move constant-product pool balances, so the ordered replay includes them alongside swaps. We unwind the sequence to the start of the period and compare the reconstructed opening balance with an independently observed reserve state,

$$\left\| \left(\mathbf{R}_{p,1}^{\text{pre}} - \mathbf{R}_{p,\text{prev}}^{\text{end}} \right) \oslash \mathbf{R}_{p,\text{prev}}^{\text{end}} \right\|_{\infty} < \tau,$$

Agreement establishes that the event history reproduces the observed reserves for the tested period. The state-dependent sample omits a period when a required event family is incomplete or the balances do not reconcile. Coverage is reported by exchange and year because missing state can be concentrated in actively managed or newly launched pools.

The same logic applies to weighted pools. We treat `poolSnapshots.amounts` as the closing balance and unwind joins, exits, and swaps in their recorded order. Table 22 compares this reconstruction with two alternatives that either hold the closing balance fixed or treat it as the opening state.

The Balancer replay includes joins and exits as well as swaps. Omitting those events assigns intraday liquidity changes to the pricing rule and changes the pool states available for the price comparison.

D Restricting hypothetical trades to observed size–depth combinations

Component-validation samples contain executed swaps and therefore pools deep enough to serve the observed trade. Hypothetical routes can demand much more output relative to pool depth, where the observed errors no longer describe pricing accuracy. We include a hypothetical leg in the cost comparison only when its own price impact is at most 5%, calculated under the pricing rule of the pool family it would traverse.

For constant-product swaps, input relative to reserves provides a transparent size–depth com-

Table 23: Size relative to depth in the validation population

Day	Swaps	Median	90th pct.	99th pct.	99.9th pct.	Maximum
2020-05-05	2	1.00%	1.00%	1.00%	1.00%	1.00%
2021-02-10	141,804	0.19%	2.00%	8.68%	23.17%	99.83%
2021-11-18	107,059	0.10%	2.24%	13.15%	71.64%	100.00%
2022-08-26	79,180	0.48%	3.66%	18.44%	94.13%	100.00%
2023-06-03	159,551	0.85%	4.93%	19.36%	96.59%	100.00%
2024-03-10	215,060	0.41%	3.61%	15.11%	57.81%	100.00%
2024-12-16	132,508	0.31%	3.05%	13.25%	71.93%	100.00%
2025-09-23	97,106	0.24%	2.51%	19.92%	99.11%	100.00%
Pooled	932,270	0.34%	3.29%	14.86%	78.33%	100.00%

Note: Trade size as a percentage of the input-side reserve of the pool traversed, over the constant-product swaps the quoters were validated against. The sample contains eight dates from 2020-05-05 through 2025-09-23.

parison,


$$\iota_{i,p,e^-} = \frac{\Delta_i}{R_{i,p,e^-}},$$

the input amount as a fraction of the input-side reserve. Table 23 reports its distribution in the historical validation sample. The statistic helps calibrate what constitutes a large trade in a constant-product pool, but it does not replace the common own-price-impact rule for concentrated-liquidity, StableSwap, or weighted pools.

In the Pooled row of Table 23, five per cent of input reserves sits between the 90th percentile of 3.29% and the 99th of 14.86%; the same cutoff lies between those columns on seven of the eight date rows. Figure 4 shows the daily distributions. This calculation locates the cutoff in the empirical size distribution. The empirical rule remains a 5% own-price-impact cutoff computed separately under each pool family’s pricing function.

Equation (11) applies the same 5% own-price-impact ceiling to every family. The component exercises establish pricing accuracy for realised swaps within their stated samples; they do not measure quote error on unexecuted routes.

Figure 4: Constant-product trade size relative to reserves



Note: The statistic is a leg’s input amount as a fraction of the input-side reserve in the constant-product pool it traverses. Points mark the 90th and 99th percentiles; the dashed line marks 5% of reserves. The two percentiles coincide on 2020-05-05. The sample contains 932,270 realised swaps on eight dates from 2020-05-05 through 2025-09-23.

E Coverage of the price comparison

Omitting an exchange can change the best direct route, the best intermediated route, or both. Table 24 first describes the volume recorded in nine local source series, including Uniswap v1 as a separate architecture sample and the available Fluid extracts. Volume in the Graph-indexed series is reconstructed from individual swaps using the smaller dollar value of the two legs. This avoids double counting in two series and removes implausible oracle values in thin pools: one Curve pool reports a \$456m day whose swaps sum to \$8, and a constant-product atomic pair with \$255 of reserves reports \$1.36bn. Physical bounds on volume relative to reserves exclude between 0.10% and 0.24% of Graph-indexed pool-days per year. Uniswap v1 is excluded from the route-cost comparison because its forced-routing architecture is studied separately. Fluid supplies volume on only 40 of the 400 sampled dates and no historical TVL field. The table consequently reports source coverage. Complete market shares require a common historical state series for every exchange.

The Uniswap v4 pricing model covers static-fee pools without hooks. Pool-identity reconstruction leaves zero swaps with incomplete or invalid immutable pool attributes across 29,005,863 rows. Pools in this group account for 84.02% of swaps and 99.05% of reported value; the corresponding 2026 shares are 79.89% and 95.57%. Other pools use hooks, dynamic fees, or both, whose transaction-level cash flows are not identified from the raw event. These shares describe the reach of the pricing model. They do not establish the availability of every pre-transaction state.

Curve carries between 9.68% and 18.70% of observed source volume in every year it exists and ranks second in four years, making its coverage economically important. The StableSwap model fits

Table 24: Share of observed decentralised-exchange volume, by source and year

Year	Uni V1	Uni V2	Uni V3	Uni V4	Sushi V2	Sushi V3	Curve	Balancer	Fluid
2020	2.09	75.89	0.00	0.00	10.89	0.00	11.12	0.00	0.00
2021	0.02	31.27	37.97	0.00	17.31	0.00	12.25	1.18	0.00
2022	0.01	7.01	64.07	0.00	4.35	0.00	18.70	5.87	0.00
2023	0.01	9.26	66.81	0.00	1.21	0.03	13.88	8.80	0.00
2024	0.00	11.63	69.06	0.00	0.49	0.01	12.57	5.86	0.38
2025	0.01	4.21	53.10	19.53	0.28	0.01	9.68	1.49	11.70
2026	0.01	2.27	46.02	31.94	0.11	0.16	12.61	0.15	6.73
Pooled	0.05	14.58	54.54	5.30	5.73	0.02	13.37	3.82	2.61

Note: Percentage shares of volume recorded in nine local source series. The weekly calendar contains 400 dates; Fluid is observed on 40 (6 in 2024, 25 in 2025, and 9 in 2026) and supplies no TVL field. Its entries are partial observed-volume contributions, not complete market shares or turnover estimates. The table displays 2020–2026 and pools volume over those years. Uniswap v1 is shown as the separate forced-routing architecture sample and is excluded from the route-cost opportunity set.

621 pool-days and excludes 588 whose 90th-percentile fitting error exceeds 1%. Across 28 sampled dates from 2020-02-11 to 2026-04-28, these exclusions account for 34.2% of the Curve volume assessed, or \$1.83bn of \$5.36bn. Another 2.75% of Curve volume occurs in pool-days with too few trades to estimate and evaluate the amplification coefficient separately. Lowering the minimum from eight trades to four changes annual excluded shares by at most 1.3 pp. Table 25 shows where these exclusions occur.

The exclusions are concentrated in routes involving the native asset. In Panel B of Table 25, the excluded share of native-leg volume exceeds the stable-leg share by at least 33 pp in every sample year and by 54 pp in the 2023 row. Tricrypto pools pairing WBTC and WETH with a stablecoin account for \$818m, or 44.6%, of the \$1.83bn excluded, including seven of the twelve largest excluded pool-days. These pools use CryptoSwap, so the missing coverage is concentrated where Curve supplies depth between ETH and stablecoins.

SushiSwap v3 is omitted from the route-cost comparison because its historical records lack the pool’s root price, in-range liquidity, and tick spacing. Its reserves are distributed across ticks, and the balance and weight fields exposed by the common schema do not describe that invariant. The exchange accounts for 0.016% of volume among the seven exchanges used for route costs over the full sample and 0.176% in 2026. Across fourteen sampled days, only 4.1% of its volume occurs on token atomic pairs unavailable on another included exchange, never more than \$37,609 on one day.

Balancer accounts for 3.9% of measured venue volume over the full sample and 8.8% at its 2023 peak. Table 22 validates selected weighted pools. Linear and boosted pools remain outside the quoter. Of 29 sampled composable-stable pools, 23 reduce to two external tokens after removing the pool’s own share token, but pricing them also requires historical rate-provider states absent from the source. The priced Balancer sample is therefore confined to weighted pools whose historical state can be reconstructed and whose realised trades are reproduced at one of the tiers of Equation (18).

Table 25: Curve volume outside the StableSwap comparison

<i>Panel A: composition of the exclusion</i>					
Leg served	Pool-days	Volume	Of excluded	Of priced	Of that leg
Native	374	\$1.038bn	56.6%	15.7%	65.2%
Stable	119	\$0.752bn	41.0%	79.9%	21.1%
Other volatile	95	\$0.045bn	2.4%	4.4%	22.2%
<i>Panel B: excluded share of each leg’s volume, by year</i>					
Year	Native leg		Stable leg	Other volatile	
2021	85.49%		46.91%	15.23%	
2022	68.82%		30.21%	50.77%	
2023	59.11%		5.51%	56.60%	
2024	48.24%		0.94%	14.38%	
2025	47.47%		13.96%	10.52%	
2026	60.92%		9.04%	20.08%	

Note: A pool serves the *native* leg when it contains ETH or a staking derivative, the *stable* leg when every token is stable, including baskets and interest-bearing wrappers, and the *other volatile* leg otherwise. Stablecoin base-pool tokens are classified with the underlying stable assets. The sample contains 28 measured Curve pool-days from 11 February 2020 through 28 April 2026.

F Round-trip exclusion

A route whose first input token equals its last output token,

$$\mathbf{1}\{k_0 = k_{m(r)}\},$$

does not represent a conversion between distinct endpoint assets. We exclude these routes from the dominance measures. The excluded population can contain cyclic arbitrage and other self-returning activity; the endpoint rule does not classify every route as arbitrage or wash trading. Multi-trade arbitrage can be bundled atomically (Daian et al., 2020), but identifying trader intent would require additional evidence.

Table 26 reports the distribution across sampled days; Figure 5 shows the day-level observations.


Across the quantile rows of Table 26, the Share by value column exceeds Share by count by a factor between 1.7 and 3.5, consistent with a population in which large self-returning routes carry disproportionate value. Round trips appear on every sampled day; their minimum count share is 4.45%. The paper excludes them before forming both count- and value-weighted dominance measures. Figure 5 plots their count and value shares among multi-leg routes and identifies the date of the maximum.

Table 26: Round-trip share of multi-leg routing, across 79 sampled days

Quantile across days	Share by count	Share by value
Minimum	4.45%	7.40%
25th percentile	9.60%	17.73%
Median	12.70%	21.72%
75th percentile	17.04%	29.56%
90th percentile	20.03%	46.11%
Maximum	25.93%	91.31%
2020 median	11.77%	14.08%
2021 median	13.21%	25.35%
2022 median	16.72%	21.89%
2023 median	8.51%	17.71%
2024 median	9.60%	21.13%
2025 median	16.75%	36.08%
2026 median	20.80%	20.32%

Note: Round trips as a share of multi-leg routes. The sample contains 79 dates from 1 June 2020 through 27 April 2026 and 13,160,047 routes, of which 2,488,399 are multi-leg and 351,553 are round trips. Dates with fewer than 200 multi-leg routes are omitted.

Figure 5: Round-trip contamination of multi-leg routing, day by day



Note: Each mark is one sampled date. *Share by value* weights round-trip routes by their notional; *share by count* gives each route equal weight. Both shares use multi-leg routes as the denominator. Dashed and dotted lines mark the respective medians, and the highest point is labelled. Dates with fewer than 200 multi-leg routes are omitted. The sample contains 79 dates from 1 June 2020 through 27 April 2026 and 2,488,399 multi-leg routes.

References

- Adams, H., N. Zinsmeister, M. Salem, R. Keefer, and D. Robinson (2021). Uniswap v3 core. Technical whitepaper.
- Amiti, M., O. Itskhoki, and J. Konings (2022). Dominant currencies: How firms choose currency invoicing and why it matters. *Quarterly Journal of Economics* 137(3), 1435–1493.
- Anadu, K., P. Azar, M. Cipriani, T. M. Eisenbach, C. Huang, M. Landoni, G. La Spada, M. Macchiavelli, A. Malfroy-Camine, and J. C. Wang (2025). Runs and flights to safety: Are stablecoins the new money market funds? Technical Report Supervisory Research and Analysis Working Paper SRA 23-02, Federal Reserve Bank of Boston. June 2025 revision; first circulated in 2023.
- Bank for International Settlements Innovation Hub (2024). Nexus: Foreign exchange provision, key points. Project Nexus technical documentation. Accessed 20 August 2026.
- Bank for International Settlements Innovation Hub (2025). Project rialto: Improving instant cross-border payments using central bank money settlement. Updated 10 December 2025.
- Caparros, B., A. Chaudhary, and O. Klein (2024). Blockchain scaling and liquidity concentration on decentralized exchanges. Working paper, Warwick Business School and Gillmore Centre for Financial Technology.
- Catalini, C., A. de Gortari, and N. Shah (2022). Some simple economics of stablecoins. *Annual Review of Financial Economics* 14, 117–135.
- Chen, D. and D. Duffie (2021). Market fragmentation. *American Economic Review* 111(7), 2247–2274.
- Daian, P., S. Goldfeder, T. Kell, Y. Li, X. Zhao, I. Bentov, L. Breidenbach, and A. Juels (2020). Flash boys 2.0: Frontrunning in decentralized exchanges, miner extractable value, and consensus instability. In *2020 IEEE Symposium on Security and Privacy*, pp. 910–927.
- Dowd, K. and D. Greenaway (1993). Currency competition, network externalities and switching costs: Towards an alternative view of optimum currency areas. *The Economic Journal* 103(420), 1180–1189.
- Eren, E. and S. Malamud (2022). Dominant currency debt. *Journal of Financial Economics* 144(2), 571–589.
- Flandreau, M. and C. Jobst (2009). The empirics of international currencies: Network externalities, history and persistence. *The Economic Journal* 119(537), 643–664.
- Gopinath, G., E. Boz, C. Casas, F. J. Díez, P.-O. Gourinchas, and M. Plagborg-Møller (2020). Dominant currency paradigm. *American Economic Review* 110(3), 677–719.

- Gopinath, G. and J. C. Stein (2021). Banking, trade, and the making of a dominant currency. *Quarterly Journal of Economics* 136(2), 783–830.
- Gorton, G. B. and J. Y. Zhang (2023). Taming wildcat stablecoins. *University of Chicago Law Review* 90(3), 909–971.
- Huang, W., A. Ranaldo, A. Schrimpf, and F. Somogyi (2025). Constrained liquidity provision in currency markets. *Journal of Financial Economics* 167, 104028.
- Kiyotaki, N. and R. Wright (1989, August). On money as a medium of exchange. *Journal of Political Economy* 97(4), 927–954.
- Krugman, P. (1980). Vehicle currencies and the structure of international exchange. *Journal of Money, Credit and Banking* 12(3), 513–526.
- Lehar, A. and C. A. Parlour (2025). Decentralized exchange: The uniswap automated market maker. *Journal of Finance* 80(1), 321–374.
- Lyons, R. K. and G. Viswanath-Natraj (2023). What keeps stablecoins stable? *Journal of International Money and Finance* 131, 102777.
- Makarov, I. and A. Schoar (2020). Trading and arbitrage in cryptocurrency markets. *Journal of Financial Economics* 135(2), 293–319.
- Milionis, J., C. C. Moallemi, T. Roughgarden, and A. L. Zhang (2022). Automated market making and loss-versus-rebalancing.
- Mukhin, D. (2022). An equilibrium model of the international price system. *American Economic Review* 112(2), 650–688.
- Schär, F. (2021). Decentralized finance: On blockchain- and smart contract-based financial markets. *Federal Reserve Bank of St. Louis Review* 103(2), 153–174.
- Somogyi, F. (2026). Dollar dominance in FX trading. *Management Science*.
- Xi, W. and C. C. Moallemi (2026). Quantifying sub-optimality in routing for automated market makers. Accepted at the 5th Workshop on Decentralized Finance.